

Zoo: A framework for the verification of concurrent OCAML 5 programs using separation logic

ANONYMOUS AUTHOR(S)

The release of OCAML 5, which introduced parallelism in the OCAML runtime, drove the need for safe and efficient concurrent data structures. New libraries like [Saturn](#) address this need. This is an opportunity to apply and further state-of-the-art program verification techniques.

We present Zoo, a framework for verifying fine-grained concurrent OCAML 5 algorithms. Following a pragmatic approach, we defined a limited but sufficient fragment of the language to faithfully express these algorithms: ZOO_{LANG}. We formalized its semantics carefully via a deep embedding in the [Rocq](#) proof assistant, uncovering subtle aspects of physical equality. We provide a tool to translate source OCAML programs into ZOO_{LANG} syntax embedded inside [Rocq](#), where they can be specified and verified using the [Iris](#) concurrent separation logic. To illustrate the applicability of Zoo, we verified a subset of the standard library and a collection of fine-grained concurrent data structures from the [Saturn](#) and [Eio](#) libraries.

In the process, we also extended OCAML to more efficiently express certain concurrent programs.

1 INTRODUCTION

OCaml 5.0 was released on December 15th 2022, the first version of the OCaml programming language to support parallel execution of OCAML threads by merging the MULTICORE OCAML runtime [Sivaramakrishnan, Dolan, White, Jaffer, Kelly, Sahoo, Parimala, Dhiman and Madhavapeddy 2020]. It provided basic support in the language runtime to start and stop coarse-grained threads called “domains”, and support for strongly sequential atomic references in the standard library. The third-party library `domainslib` offered a simple scheduler for a pool of tasks, used to benchmark the parallel runtime. A world of parallel and concurrent software was waiting to be invented.

Shared-memory concurrency is a difficult programming domain, and existing ecosystems (C++, Java, Haskell, Rust, Go...) took decades to evolve comprehensive libraries of concurrent abstractions and data structures. In the last couple years, a handful of contributors to the OCaml ecosystem have been implementing basic libraries for concurrent and parallel programs in OCAML, in particular [Saturn](#) [Karvonen and Morel 2025], a library of lock-free thread-safe data structures (stacks, queues, a work-stealing dequeue, a skip list, a bag and a hash table), [Eio](#) [Madhavapeddy and Leonard 2025], a library of asynchronous IO and structured concurrency, and [Kcas](#) [Karvonen 2025a], a library offering a software-transactional-memory abstraction.

Concurrent algorithms and data structures are extremely difficult to reason about. Their implementations tend to be fairly short, a few dozens of lines. There is only a handful of experts able to write such code, and many potential users. They are difficult to test comprehensively. These characteristics make them ideally suited for mechanized program verification.

We embarked on a mission to mechanize correctness proofs of OCAML concurrent algorithms and data structures as they are being written, in contact with their authors, rather than years later. In the process, we not only gained confidence in these complex new building blocks, but we also improved the OCAML language and its verification ecosystem.

OCAML language features. When studying the new codebases of concurrent and parallel data structures, we found a variety of unsafe idioms, working around expressivity or performance limitations with the OCAML language support for lock-free concurrent data structures. In particular, the support for *atomic references* in the OCAML library proved inadequate, as idiomatic concurrent data-structures need the more expressive feature of *atomic record fields*. We designed an extension of OCAML with atomic record fields, implemented it as an experimental compiler variant, and

succeeded in getting it integrated in the upstream OCAML compiler: it should be available as part of OCaml 5.4, which is not yet released at the time of writing.

Verification tools for concurrent programs. The state-of-the-art approach for mechanized verification of fine-grained concurrent algorithms is to use [IRIS](#) [[Jung, Krebbers, Jourdan, Bizjak, Birkedal and Dreyer 2018](#)], a state-of-the-art mechanized *higher-order* concurrent separation logic with *user-defined ghost state*. Its expressivity allows to precisely capture delicate invariants, and to reason about the linearization points of fine-grained concurrent algorithms (including external [[Vindum, Frumin and Birkedal 2022](#)] and future-dependent [[Chang, Jung, Sharma, Tassarotti, Kaashoek and Zeldovich 2023](#); [Jung, Lepigre, Parthasarathy, Rapoport, Timany, Dreyer and Jacobs 2020](#); [Vindum and Birkedal 2021](#)] linearization). [IRIS](#) provides a generic mechanism to define programming languages and program logics for them. Much of the existing [IRIS](#) concurrent verification work has been performed in [HEAPLANG](#), the exemplar [IRIS](#) language; a concurrent, imperative, untyped, call-by-value functional language.

To the best of our knowledge, it is currently the closest language to OCAML 5 in the [IRIS](#) ecosystem — we review the existing frameworks in [Section 12](#). We started our verification effort in [HEAPLANG](#), but it eventually proved impractical to verify realistic OCAML libraries. Indeed, it lacks basic abstractions such as algebraic data types (tuples, mutable and immutable records, variants) and mutually recursive functions. Verifying OCAML programs in [HEAPLANG](#) requires difficult translation choices and introduces various encodings, to the point that the relation between the source and verified programs can become difficult to maintain and reason about. It also has very few standard data structures that can be directly reused. These limitations are well-known in the [IRIS](#) community.

We created a new [IRIS](#) language, [ZooLANG](#), that can better express concurrent OCAML programs. Its feature set grew over time as we applied it to more verification scenarios, and we now believe that it allows practical verification of fine-grained concurrent OCAML 5 programs — including the use of our atomic record fields which were co-designed with [ZooLANG](#). We were influenced by the [PERENNIAL](#) framework [[Chajed, Tassarotti, Kaashoek and Zeldovich 2019](#)], which achieved similar goals for the Go language with a focus on crash-safety. As in [PERENNIAL](#), we also provide a translator from (a subset of) OCAML to [ZooLANG](#): `ocaml2zoo`. We start from OCAML code and call our translator to obtain a deep [ZooLANG](#) embedding inside [Rocq](#); we can use lightweight annotations to guide the translation. Inside [Rocq](#) we define specifications using [IRIS](#), and prove them correct with respect to the [ZooLANG](#) version, which is syntactically very close to the original OCAML source. We call the resulting framework [Zoo](#).

One notable current limitation of [ZooLANG](#) is that it assumes a sequentially-consistent memory model, whereas OCAML offers a weaker memory model [[Dolan, Sivaramakrishnan and Mahavapeddy 2018](#)]. We made the choice to ensure that we supported practical verification in a sequentially-consistent setting first; in the future we plan to equip [ZooLANG](#) with the OCAML memory model as formalized in [Cosmo](#) [[Mével, Jourdan and Pottier 2020](#)]. We discuss the impact of this difference in ??.

Specified OCAML semantics. Our [IRIS](#) mechanization of [ZooLANG](#) defines an operational semantics and a corresponding program logic. Our users on the other hand run their program through the standard OCAML implementation, which is not verified and does not have a precise formal specification. To bridge this formal-informal gap as well as reasonably possible, we carefully audit our [ZooLANG](#) semantics to ensure that they coincide with OCAML's.

In doing so we discovered a hole in state-of-the-art language semantics for program verification (not just for OCAML), which is the treatment of *physical equality* (pointer quality). Physical equality is typically exposed to language users as an efficient but under-specified equality check, as the physical identity of objects may or may not be preserved by various compiler transformations. It

is an essential aspect of concurrent programs, as it underlies the semantics of important atomic instructions such as `compare_and_set`. We found that the current informal semantics in OCAML is incomplete, it does not allow to reason on programs that use structured data which mix mutable and immutable constructors. Existing formalizations of physical equality in verification frameworks typically restrict it to primitive datatypes, but idiomatic concurrent programs do not fit within this restriction. We propose a precise specification of physical equality in Zoo that scales to the verification of all the concurrent programs we encountered.

Worse, our discussions with the maintainers of the OCAML implementation showed that implementors guarantee weaker properties of physical equalities than users assume, in particular they may allow *unsharing*, which makes some existing concurrent programs incorrect. We propose a small new language feature for OCAML, per-constructor unsharing control, which we also integrate in our ZOO_{LANG} translation, to fix affected programs and verify them. Finally, we discussed these subtleties with authors who axiomatize physical equality within Rocq for the purpose of efficient extraction, and we found out that some subtleties we discovered could translate into incorrectness in their axiomatization, requiring careful restrictions.

Verification results. We verified a small library for ZOO_{LANG}, typically a subset of the OCAML standard library. It can serve as building blocks to define our concurrent data structures. (The lack of such a reusable standard library is a current limitation of [HEAP_{LANG}](#).) We verified a specific component of the [Eio](#) library, whose author Thomas Leonard had pointed to us as being delicate to reason about and worth mechanizing. Finally, we verified a large subset of the [Saturn](#) library: stacks, queues (list-based and stack-based), and finally the Chase-Lev work-stealing queue [[Chase and Lev 2005](#)]. The [Saturn](#) implementation of these lock-free data structures are used by the concurrent schedulers proposed in the OCAML 5 library ecosystem, notably `domainslib` and `picos`. (The main [Saturn](#) concurrent structures missing from our verification are a skip-list and a hashtable.) Several of these data structures contained verification challenges, which we will describe in the relevant section.

Contributions. In summary, we claim the following contributions:

- (1) Zoo, a practical program verification framework aimed at concurrent OCAML programs, mechanized in Rocq. The language ZOO_{LANG} comes with a program logic expressed in the [IRIS](#) concurrent separation logic. A translator `ocaml2zoo` generates Rocq embeddings from source OCAML programs, and works well with OCAML tooling (dune support).
- (2) The verification (in a sequentially-consistent model) of important structures coming from [Saturn](#), the OCAML 5 library of lock-free data structures. In particular we present a precise concurrent invariant for the Chase-Lev work-stealing queue, which gives stronger specifications than its previous formalizations.
Gabriel{It would be nice to emphasize here a “proof technique” for concurrent verification that is novel in Clément’s work. (Some novel usage setup for prophecy variables?)}
- (3) The extension of OCAML with atomic record fields, which after significant design, implementation and discussion work have now been integrated into upstream OCAML.
- (4) The identification of blind spots in existing specifications of *physical equality* and a new specification that is precise enough to reason about `compare-and-set` in the various programs we considered.

In the process we identified a potential bug in existing OCAML programs related to *unsharing*, and we propose a small language extension to let users selectively disable unsharing.

```

148   type 'a t = 'a list Atomic.t      Definition stack_create : val :=
149   let create () = Atomic.make []     fun: <> => ref [].
150
151   let rec push t v =                 Definition stack_push : val :=
152     let old = Atomic.get t in        rec: "push" "t" "v" =>
153     let new_ = v :: old in           let: "old" := !"t" in
154     if not (Atomic.compare_and_set   let: "new_" := "v" :: "old" in
155       t old new_)                   if: ~ CAS "t".[contents] "old" "new_"
156     then (                           then (
157       Domain.cpu_relax () ;          domain_cpu_relax () ;;
158       push t v                       "push" "t" "v"
159     )                                ).
160
161   let rec pop t =                    Definition stack_pop : val :=
162     match Atomic.get t with          rec: "pop" "t" =>
163     | [] -> None                     match: !"t" with
164     | v :: new_ as old ->            | [] => $None
165       if Atomic.compare_and_set     | ("v" :: "new_") as "old" =>
166         t old new_                  if: CAS "t".[contents] "old" "new_"
167       then Some v                   then 'Some( "v" )
168       else (                         else (
169         Domain.cpu_relax () ;        domain_yield () ;;
170         pop t                        "pop" "t"
171       )                              )
172                                     end.

```

Fig. 1. A concurrent stack in OCAML and its automatic ZooLANG translation

2 ZOO IN PRACTICE

The core of Zoo is ZooLANG: a concurrent, imperative, untyped, functional programming language fully formalized in Rocq. Its semantics has been designed to match OCAML.

ZooLANG comes with a program logic based on Iris: reasoning rules expressed in separation logic (including rules for the different constructs of the language) along with Rocq tactics that integrate into the Iris proof mode [Krebbers, Jourdan, Jung, Tassarotti, Kaiser, Timany, Charguéraud and Dreyer 2018]. In addition, it supports DIAFRAME [Mulder, Krebbers and Geuvers 2022], enabling proof automation.

2.1 Translation from OCAML to ZooLANG

While ZooLANG lives in Rocq, we want to verify OCAML programs. To connect them we provide the tool `ocaml2zoo` to translate OCAML source files¹ into Rocq files containing ZooLANG code. This tool can process entire dune projects, and support several libraries provided together or as dependencies of the project.

The supported OCAML fragment includes: tuples, variants, records and inline records, shallow **match**, atomic record fields, unboxed types, toplevel mutually recursive functions.

¹Actually, `ocaml2zoo` processes binary annotation files (.cmt files).

Rocq term	t	
constructor	C	
projection	$proj$	
record field	fld	
identifier	s, f	$\in \text{String}$
integer	n	$\in \mathbb{Z}$
boolean	b	$\in \mathbb{B}$
binder	x	$::= \langle \rangle \mid s$
unary operator	\oplus	$::= \sim \mid -$
binary operator	\otimes	$::= + \mid - \mid * \mid \text{'quot'} \mid \text{'rem'} \mid \text{'land'} \mid \text{'lor'} \mid \text{'lsl'} \mid \text{'lsr'}$ $\mid \leq \mid < \mid \geq \mid > \mid = \mid \neq \mid == \mid !=$ $\mid \text{and} \mid \text{or}$
expression	e	$::= t \mid s \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e \mid e_1 e_2$ $\mid \text{let: } x := e_1 \text{ in } e_2 \mid e_1 ; ; e_2$ $\mid \text{let: } f x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{letrec: } f x_1 \dots x_n := e_1 \text{ in } e_2$ $\mid \text{let: } \langle C x_1 \dots x_n := e_1 \text{ in } e_2 \mid \text{let: } x_1, \dots, x_n := e_1 \text{ in } e_2$ $\mid \oplus e \mid e_1 \otimes e_2$ $\mid \text{if: } e_0 \text{ then } e_1 \text{ (else } e_2) ?$ $\mid \text{for: } x := e_1 \text{ to } e_2 \text{ begin } e_3 \text{ end}$ $\mid \S C \mid \langle C (e_1, \dots, e_n) \mid (e_1, \dots, e_n) \mid e. \langle proj \rangle$ $\mid [] \mid e_1 :: e_2$ $\mid \langle C \{e_1, \dots, e_n\} \mid \{e_1, \dots, e_n\} \mid e. \{fld\} \mid e_1 \leftarrow \{fld\} e_2$ $\mid \text{ref } e \mid !e \mid e_1 \leftarrow e_2$ $\mid \text{match: } e_0 \text{ with } br_1 \mid \dots \mid br_n \mid _ \text{ (as } s) ? \Rightarrow e) ? \text{ end}$ $\mid e. [fld] \mid \text{Xchg } e_1 e_2 \mid \text{CAS } e_1 e_2 e_3 \mid \text{FAA } e_1 e_2$ $\mid \text{Proph} \mid \text{Resolve } e_0 e_1 e_2$
branch	br	$::= C (x_1 \dots x_n) ? \text{ (as } s) ? \Rightarrow e$ $\mid [] \text{ (as } s) ? \Rightarrow e \mid x_1 :: x_2 \text{ (as } s) ? \Rightarrow e$
toplevel value	v	$::= t \mid \#n \mid \#b$ $\mid \text{fun: } x_1 \dots x_n \Rightarrow e \mid \text{rec: } f x_1 \dots x_n \Rightarrow e$ $\mid \S C \mid \langle C (v_1, \dots, v_n) \mid (v_1, \dots, v_n)$ $\mid [] \mid v_1 :: v_2$

Fig. 2. ZooLANG syntax (omitting mutually recursive toplevel functions)

In Figure 1 we include the OCAML implementation of a simple lock-free concurrent stack [Treiber 1986], and its automatic translation to ZooLANG, demonstrating that readability is preserved. Readability is important as users constantly see program fragments during interactive verification.

2.2 The full language

The ZooLANG syntax is given in Figure 2², omitting mutually recursive toplevel functions that are treated specially. Expressions include standard constructs like booleans, integers, anonymous functions (that may be recursive), applications, **let** bindings, sequence, unary and binary operators, conditionals, **for** loops, tuples. In any expression, one can refer to a Rocq term representing a

²More precisely, it is the syntax of the surface language, including Rocq notations.

ZOO_{LANG} value (of type `val`) using its [Rocq](#) identifier. ZOO_{LANG} is deeply embedded: variables (bound by functions and `let`) are quoted as strings.

Data constructors (immutable memory blocks) are supported: `$C` represents a constant constructor (e.g. `$None`), and `'C (e1, ..., en)` represents a non-constant constructor (e.g. `'Some(e)`). Projections of the form `e.<proj>` (include on tuples: `(x, y).<1>`) can be used to obtain a specific component of a tuple or data constructor. ZOO_{LANG} supports shallow pattern matching (patterns cannot be nested) on data constructors with an optional fallback case.

Mutable memory blocks are constructed using either the untagged record syntax `{e1, ..., en}` or the tagged record syntax `'C {e1, ..., en}`. Reading a record field can be performed using `e.{fld}` and writing to a record field using `e1 <- {fld} e2`. Pattern matching can also be used on mutable tagged blocks provided that cases do not bind anything — in other words, only the tag is examined, no memory access is performed. Mutable references are supported: `ref e` creates a reference, `!e` reads a reference and `e1 <- e2` writes into a reference. There is no built-in syntax for arrays, they are supported through the [Array](#) standard library module (e.g. `array_make`).

Note that ZOO_{LANG} follows OCAML in sometimes eschewing orthogonality to provide more compact memory representations: constructors are n -ary instead of taking a tuple as parameter, and the tagged record syntax is distinct from a constructor taking a mutable record as parameter. In each case the simplifying encoding would introduce an extra indirection in memory, which is absent from the ZOO_{LANG} semantics. Performance-conscious experts care about these representation choices, and we care about faithfully modeling their programs.

Parallelism is mainly supported through the [Domain](#) standard library module, including domain-local storage. Atomic operations are provided as built-in constructs (`Xchg`, `CAS`, `FAA`; see [Section 3.3](#)).

The `Proph` and `Resolve` constructs model *prophecy variables* [Jung, Lepigre, Parthasarathy, Rapoport, Timany, Dreyer and Jacobs 2020], see [Section 3.4](#).

2.3 Specifications and proofs

Once the translation to ZOO_{LANG} is done, the user can write specifications and prove them in [IRIS](#). For instance, the specification of the `stack_push` function from [Figure 1](#) could be:

```
Lemma stack_push_spec t ι v :
  <<< stack_inv t ι
    | ∀ vs, stack_model t vs >>>
    stack_push t v @ ↑ι
  <<< stack_model t (v :: vs)
    | RET (); True >>>.
Proof. ... Qed.
```

Here, we use a *logically atomic specification* [da Rocha Pinto, Dinsdale-Young and Gardner 2014], which has been proven [Birkedal, Dinsdale-Young, Guéneau, Jaber, Svendsen and Tzevelekos 2021] to be equivalent to *linearizability* [Herlihy and Wing 1990] in sequentially consistent memory models.

Similarly to [Hoare triples](#), the specification is formed of a precondition and a postcondition, represented in angle brackets. But each is split in two parts, a *public* or *atomic* condition, and a *private* condition. Following standard [IRIS](#) notations, the private conditions are on the outside (first line of the precondition, last line of the postcondition) and the atomic conditions are inside.

For this particular operation, the private postcondition is trivial. The private precondition `stack_inv t` is the stack invariant. Intuitively, it asserts that `t` is a valid concurrent stack. More precisely, it defines a concurrent protocol that `t` must respect at all times.

The atomic pre- and post-conditions specify the linearization point of the operation: during the execution of `stack_push`, the abstract state of the stack held by `stack_model` is atomically updated from vs to $v :: vs$ when v is atomically pushed at the top of the stack.

3 ZOO FEATURES

In this section, we review the salient features of Zoo, which we found lacking when we attempted to use `HEAPLANG` to verify real-world OCAML programs. Providing a better `IRIS` language than `HEAPLANG` for this problem domain is a contribution of our work; others are welcome to use `ZOO LANG` for their own verification effort, or to reuse our designs by integrating specific features into their own `IRIS` language.

3.1 Algebraic data types

`ZOO LANG` is an untyped language but, to write interesting programs, it is convenient to work with abstractions like algebraic data types. To simulate tuples, variants and records, we designed a machinery to define projections, constructors and record fields.

For example, one may define a list-like type with:

Notation `"Nil"` := (in_type "t" 0) (in custom zoo_tag).

Notation `"Cons"` := (in_type "t" 1) (in custom zoo_tag).

Users do not need to write this incantation directly, as they are generated by `ocaml2zoo` from the OCAML type declarations. Suffice it to say that it introduces the two tags in the `zoo_tag` custom entry, on which the notations for data constructors rely. The `in_type` term is needed to distinguish the tags of distinct data types; crucially, it cannot be simplified away by `Rocq`, as this could lead to confusion during the reduction of expressions.

One may then directly use the tags `Nil` and `Cons` in data constructors:

Definition `map` : val :=
 rec: "map" "fn" "t" =>
 match: "t" with
 | Nil => \$Nil
 | Cons "x" "t" =>
 let: "y" := "fn" "x" in
 'Cons("y", "map" "fn" "t")
 end.

Similarly, one can define a record-like type with two mutable fields `f1` and `f2`:

Notation `"f1"` := (in_type "t" 0) (in custom zoo_field).

Notation `"f2"` := (in_type "t" 1) (in custom zoo_field).

Definition `swap` : val :=
 fun: "t" =>
 let: "f1" := "t".{f1} in
 "t" <-{f1} "t".{f2} ;; "t" <-{f2} "f1".

3.2 Mutually recursive functions

Zoo supports non-recursive (fun: $x_1 \dots x_n \Rightarrow e$) and recursive (rec: $f \ x_1 \dots x_n \Rightarrow e$) functions but only *oplevel* mutually recursive functions. It is non-trivial to properly handle mutual recursion: when applying a mutually recursive function, a naive approach would replace calls to sibling functions by their respective bodies, but this typically makes the resulting expression unreadable and can create proof-checking performance issues during verification. To prevent it, the mutually

recursive functions have to know one another to preserve their names during β -reduction. We simulate this using some boilerplate that can be generated by `ocaml2zoo`. For instance, one may define two mutually recursive functions `f` and `g` as follows:

```

Definition f_g := (
  recs: "f" "x" => "g" "x"
  and:  "g" "x" => "f" "x"
)%zoo_recs.

(* boilerplate *)
Definition f := ValRecs 0 f_g.
Definition g := ValRecs 1 f_g.
Instance : AsValRecs' f 0 f_g [f;g]. Proof. done. Qed.
Instance : AsValRecs' g 1 f_g [f;g]. Proof. done. Qed.

```

3.3 Fine-grained concurrent primitives

Zoo supports concurrent primitives both on atomic references (from `Atomic`) and atomic record fields (from `Atomic.Loc`³) according to the table below. The OCAML expressions listed in the left-hand column translate into the Zoo expressions in the right-hand column. Notice that an atomic location `[%atomic.loc e.f]` (of type `_ Atomic.Loc.t`) translates directly into `e.[f]`.

OCAML	Zoo
<code>Atomic.get e</code>	<code>!e</code>
<code>Atomic.set e₁ e₂</code>	<code>e₁ <- e₂</code>
<code>Atomic.exchange e₁ e₂</code>	<code>Xchg e₁. [contents] e₂</code>
<code>Atomic.compare_and_set e₁ e₂ e₃</code>	<code>CAS e₁. [contents] e₂ e₃</code>
<code>Atomic.fetch_and_add e₁ e₂</code>	<code>FAA e₁. [contents] e₂</code>
<code>Atomic.Loc.exchange [%atomic.loc e₁.f] e₂</code>	<code>Xchg e₁. [f] e₂</code>
<code>Atomic.Loc.compare_and_set [%atomic.loc e₁.f] e₂ e₃</code>	<code>CAS e₁. [f] e₂ e₃</code>
<code>Atomic.Loc.fetch_and_add [%atomic.loc e₁.f] e₂</code>	<code>FAA e₁. [f] e₂</code>

One notable aspect of this translation is that atomic accesses (`Atomic.get` and `Atomic.set`) correspond to plain loads and stores. This is because we are working in a sequentially consistent memory model: there is no difference between atomic and non-atomic memory locations.

3.4 Prophecy variables

Lock-free algorithms exhibit complex behaviors. To tackle them, `IRIS` provides powerful mechanisms such as *prophecy variables* [Jung, Lepigre, Parthasarathy, Rapoport, Timany, Dreyer and Jacobs 2020]. Essentially, prophecy variables can be used to predict the future of the program execution and reason about it. They are key to handle *future-dependent linearization points*: linearization points that may or may not occur at a given location in the code depending on a future observation.

Zoo supports prophecy variables through the `Proph` and `Resolve` expressions — as in `HEA-PLANG`, the canonical `IRIS` language. In OCAML, these expressions correspond to `Zoo.proph` and `Zoo.resolve`, that are recognized by `ocaml2zoo`.

4 OCAML EXTENSIONS FOR FINE-GRAINED CONCURRENT PROGRAMMING

Over the course of this work, we studied efficient fine-grained concurrent OCAML programs written by experts. This revealed various limitations of OCAML in these domains, that those experts would work around using unsafe casts, often at the cost of both readability and memory-safety; and also

³The `Atomic.Loc` module is part of the `PR` that implements atomic record fields (see Section 4).

some mismatches between their mental model of the semantics of OCAML and the mental model used by the OCAML compiler authors. We worked on improving OCAML itself to reduce these work-arounds or semantic mismatches.

4.1 Atomic record fields

OCAML 5 offers a type `'a Atomic.t` of atomic references exposing sequentially-consistent atomic operations. Data races on non-atomic mutable locations has a much weaker semantics and is generally considered a programming error. For example, the Michael-Scott concurrent queue [Michael and Scott 1996] relies on a linked list structure that could be defined as follows:

```
type 'a node = Nil | Cons of { value : 'a; next : 'a node Atomic.t }
```

Performance-minded concurrency experts dislike this representation, because `'a Atomic.t` introduces an indirection in memory: it is represented as a pointer to a block containing the value of type `'a`. Instead, they use something like the following:

```
type 'a node = Nil | Cons of { mutable next: 'a node; value: 'a }
let as_atomic : 'a node -> 'a node Atomic.t option = function
| Nil -> None
| (Cons _) as record -> Some (Obj.magic record : 'a node Atomic.t)
```

Notice that the `next` field of the `Cons` constructor has been moved first in the type declaration. Because the OCAML compiler respects field-declaration order in data layout, a value `Cons { next; value }` has a similar low-level representation to a reference (atomic or not) pointing at `next`, with an extra argument. The code uses `Obj.magic` to unsafely cast this value to an atomic reference, which appears to work as intended.

`Obj.magic` is a shunned unsafe cast (the OCAML equivalent of `unsafe` or `unsafePerformIO`). It is very difficult to be confident about its usage given that it may typically violate assumptions made by the OCAML compiler and optimizer. In the example above, casting a two-fields record into a one-argument atomic reference may or may not be sound — but it gives measurable performance improvements on concurrent queue benchmarks.

It is possible to statically forbid passing `Nil` to `as_atomic` to avoid error handling, by turning `'a node` into a GADT indexed over a type-level representation of its head constructor. Examples of this pattern can be found in the `Kcas` [Karvonen 2025a] library by Vesa Karvonen. It is difficult to write correctly and use, in particular as unsafe casts can sometimes hide type-errors in the intended static discipline.

Note that this unsafe approach only works for the first field of a record, so it is not applicable to records that hold several atomic fields, such as the toplevel record storing atomic front and back pointers for the concurrent queue.

4.1.1 Our atomic fields proposal. In May-June 2024 we proposed a design for atomic record fields as an OCAML language change proposal: RFC #39⁴. Declaring a record field atomic simply requires an `[@atomic]` attribute — and could eventually become a proper keyword of the language.

```
(* re-implementation of atomic references *)
type 'a atomic_ref = { mutable contents : 'a [@atomic]; }

(* concurrent linked list *)
type 'a node = Nil | Cons of { value: 'a; mutable next : 'a node [@atomic]; }

(* bounded SPSC circular buffer *)
```

⁴De-anonymizing link: <https://github.com/ocaml/RFCs/pull/39>

```

442 type 'a bag =
443   { data : 'a Atomic.t array;
444     mutable front: int [@atomic];
445     mutable back: int [@atomic]; }

```

The design difficulty is to express atomic operations on atomic record fields. For example, if `buf` has type `'a bag` above, then one naturally expects the existing notation `buf.front` to perform an atomic read and `buf.front <- n` to perform an atomic write. But how would one express exchange, compare-and-set and fetch-and-add? We would like to avoid adding a new primitive language construct for each atomic operation.

Two different designs have been proposed for atomic operations on atomic record fields.

- (1) The original, “full” design in our proposal introduces a built-in type `('r, 'a) Atomic.Field.t` that denotes a field/index of type `'a` within a record of type `'r`, with a syntax extension `[%atomic.loc <field>]` to construct such field description, and atomic primitives in a module `Atomic.Field`, that need both the record value of type `'r` and the field description. For example, the standard library exposes:

```

457 val Atomic.Field.fetch_and_add :
458   'a -> ('a, int) Atomic.Field.t -> int -> int

```

and users can write:

```

460 let preincrement_front (buf : 'a bag) : int =
461   Atomic.Loc.fetch_and_add buf [%atomic.field front] 1

```

where `[%atomic.field front]` has type `('a bag, int) Atomic.Loc.t`. Internally, a value of type `('r, 'a) Atomic.Field.t` is just an integer offset locating the field within the record.

- (2) An alternative “simple” design, which was proposed by Basile Clément, introduces a built-in type `'a Atomic.Loc.t` for an atomic location that holds an element of type `'a`, with a syntax extension `[%atomic.loc <expr>.<field>]` to construct such locations. Atomic primitives operate on values of type `'a Atomic.Loc.t` and they are exposed as functions of the module `Atomic.Loc`. For example, the standard library exposes:

```

470 val Atomic.Loc.fetch_and_add :
471   int Atomic.Loc.t -> int -> int

```

and users can write:

```

473 let preincrement_front (buf : 'a bag) : int =
474   Atomic.Loc.fetch_and_add [%atomic.loc buf.front] 1

```

where `[%atomic.loc buf.front]` has type `int Atomic.Loc.t`. Internally, a value of type `'a Atomic.Loc.t` can be represented as a pair of a record and an integer offset for the desired field, and the `atomic.loc` construction builds this pair in a well-typed manner. When a primitive of the `Atomic.Loc` module is applied to an `atomic.loc` expression, the compiler can optimize away the construction of the pair — but it would happen if there was an abstraction barrier between the construction and its use.

The simple design has the strong advantage of exposing simpler types to users. The full design involves more complex types, but a simpler data representation that does not rely on specific compiler optimizations to generate efficient code, even across abstraction boundaries. As remarked by Leo White, the simpler type `'a Atomic.Loc.t` can be reconstructed as a dependent pair of a `'r` and a `('r, 'a) Atomic.Field.t`, which is expressible in OCAML as a GADT:

```

487 type 'a loc = Loc : 'r * ('r, 'a) Atomic.Field.t -> 'a loc

```

In August and September 2024 we implemented the simple design as an experimental version of the OCAML compiler. We got a full code review by Olivier Nicole. When we approached the OCAML

maintainers to integrate the feature upstream, they asked us to reconsider implementing the full design, based on feedback in this direction by Vesa Karvonen (author of [Kcas](#) and [Saturn](#)).

We implemented the full design between November 2024 and January 2025. We found that it is harder to implement in the type-checker. The extension form `[%atomic.loc buf.front]` of the simple design has typing rules that are very similar to a field access `buf.front`. On the other hand, `[%atomic.loc front]` interacts in a non-trivial way with the OCAML machinery for type-based disambiguation of record fields — several records with a field named `front` can co-exist in the typing environment.

For technical reasons, there is also a non-trivial interaction with the type-checking of inline record types (record types that are not defined by themselves but only as the argument of a sum type constructor), which currently prevents from using this approach with inline records, who are unfortunately common in efficient mutable data structures. In February and March 2025 we worked with OCAML maintainers experts to try to lift these limitations and proposed implementation changes for inline records, but we failed to reach a consensus on the best design.

In April and May, we managed to build consensus among OCAML maintainers that the simple design, with a simpler user-facing design and a simpler implementation, was a good compromise in view of the significant difficulties of the full design with inline records. Our implementation of the simple design was reviewed again, we wrote user documentation, and it was integrated upstream in May 2025, to be included in the upcoming release of OCaml 5.4.

Limitation: no support for cache contention. The type `Atomic.t` comes with a function

```
val Atomic.make_contended : 'a -> 'a Atomic.t
```

that ensures that the returned atomic reference is allocated with enough alignment and padding to sit alone on its cache line, to avoid performance issues caused by false sharing. Currently there is no such support for padding of atomic record fields (we are planning to work on this if the support for atomic fields gets merged in standard OCAML), so the less-compact atomic references remain preferable in certain scenarios.

4.2 Atomic arrays

On top of our atomic record fields, we have implemented experimental support for atomic arrays, another facility commonly requested by authors of efficient concurrent programs. Our previous example of a concurrent bag of type `'a bag` used a backing array of type `'a Atomic.t array`, which contains more indirections than may be desirable, as each array element is a pointer to a block containing the value of type `'a`, instead of storing the value of type `'a` directly in the array.

Our implementation of atomic arrays⁵ builds on top of the type `'a Atomic.Loc.t` we described in the previous section, and it relies on two new low-level primitives provided by the compiler:

```
val Atomic_array.index : 'a array -> int -> 'a Atomic.Loc.t
val Atomic_array.unsafe_index : 'a array -> int -> 'a Atomic.Loc.t
```

The function `index` takes an array and an integer index within the array, and returns an atomic location into the corresponding element after performing a bound check. `unsafe_index` omits the boundcheck — additional performance at the cost of memory-safety — and allows to express the atomic counterpart of the unsafe operations `Array.unsafe_get` and `Array.unsafe_set`. The atomic primitives of the module `Atomic.Loc` can then be used on these indices; our implementation implements a library module on top of these primitives to provide a higher-level layer to the user, with direct array operations such as:

⁵Non-anonymous link: <https://<clickable>>

```

540 val Atomic_array.exchange : 'a Atomic_array.t -> int -> 'a -> 'a
541 val Atomic_array.unsafe_exchange : 'a Atomic_array.t -> int -> 'a -> 'a

```

5 PHYSICAL EQUALITY

The notion of *physical equality* is ubiquitous in fine-grained concurrent algorithms. It appears not only in the semantics of the (`==`) operator, but also in the semantics of the `Atomic.compare_and_set` primitive, which atomically sets an atomic reference to a desired value if its current content is physically equal to an expected value. This primitive is commonly used to try committing an atomic operation in a retry loop, as in the push and pop functions of Figure 1.

5.1 Physical equality in `HEAPLANG`

In `HEAPLANG`, this primitive is provided but restricted. Indeed, its semantics is only defined if either the expected or the desired value fits in a single memory word in the `HEAPLANG` value representation: literals (booleans, integers and pointers⁶) and literal injections⁷; otherwise, the program is stuck. In practice, this restriction forces the programmer to introduce an indirection [Iris development team 2025; Jung, Lepigre, Parthasarathy, Rapoport, Timany, Dreyer and Jacobs 2020; Vindum and Birkedal 2021] to physically compare complex values, e.g. lists. Furthermore, when the semantics is defined, values are compared using their `Rocq` representations; physical equality boils down to `Rocq` equality.

5.2 Physical equality in `OCAML`

In `OCAML`, physical equality is more tricky. *Structural equality* `v1 = v2`, which we describe in Section 6, would be the preferred way of comparing values, and using physical equality `v1 == v2` is often an unintentional mistake. However, physical equality is typically much faster than structural equality, as it compiles to only one assembly instruction instead of traversing the value. Also, the `Atomic.compare_and_set` requires the comparison to be atomic, ruling out structural equality.

Physical equality is in a counter-intuitive situation: it is very simple to *implement* (in the `OCAML` compiler, or in an interpreter, etc.) but difficult to *specify* precisely. To make verification practical, we need a specification at the level of *source* `OCAML` (or `ZOO`) programs, using a high-level representation of values, as close to the source as reasonably possible, that we call *abstract* values. On the other hand, its implementation typically work with *low-level* values, and its observable behavior depends on compiler transformations that happen in-between the two abstraction levels. This difficulty can result in dangerous gaps between the programming language used to write code and the semantics used for its verification.

`ZOOLANG` has a grammar of values, and most operations are specified by defining how they compute with `ZOOLANG` values. Its definition may look as follows in `Rocq` (simplified):

```

576 Inductive literal :=
577   | Bool (b : bool)
578   | Int (n : nat)
579   | Loc (l : location)
580   | Proph (pid : prophet_id)
581   | Poison.

```

```

582
583 Inductive val :=
584

```

⁶`HEAPLANG` allows arbitrary pointer arithmetic and therefore inner pointers. This is forbidden in both `OCAML` and `ZOOLANG`, as any reachable value has to be compatible with the garbage collector.

⁷`HEAPLANG` has no primitive notion of constructor, only pairs and injections (left and right).

```

589 | Lit (lit : literal)
590 | Recs (i : nat) (recs : list (binder * binder * expr))
591 | Block (tag : nat) (vs : list val).

```

The value `'Cons(42, $Nil)` is represented in `Rocq` as `Block 1 [Lit (Int 42), Block 0 []]`. Notice that immutable blocks are represented in `Rocq` using the `Block` constructor directly, and *not* as a location (`Loc`) allocated on the heap. We use locations only for *mutable* records. We would say that our representation of `ZooLANG` values is high-level or *abstract*, as close to the surface syntax as reasonably possible. This distinction is important to make verification pleasant in practice, by reducing the number of locations and heap indirections that the programmer needs to work with during verification. A `ZooLANG` tuple is directly a tuple, etc., and this design decision of using high-level values is important to the verification experience — in addition, assuming full ownership of arguments of immutable blocks would be incorrect.

It is tempting to specify, as `HEAPLANG` does, that physical equality decides equality between abstract values. This specification makes sense for immediate values (integers, booleans, constant constructors), and for mutable records which are compared by location. But it is incorrect on immutable blocks, and `HEAPLANG` essentially does not specify its behavior on those values. Yet programmers use physical equality on immutable blocks in practice, as in our example of a Treiber stack of [Figure 1](#).

Defining physical equality as `Rocq` equality of abstract values is problematic in opposite ways:

- (1) Some distinct abstract values are physically equal in OCaml, for example `0` and `false`. Their type differ, but it possible to store them in an existential type where they can be compared for physical equality:

```

611 type any = Any : 'a -> any
612 let test1 = Any false == Any 0 (* may return true *)

```

This shows that even on immediate values, specifying physical equality as equality of abstract values is convenient but incorrect in practice.

- (2) A deeper problem is that some definitionally equal abstract values may be physically distinct. Consider for example the case of *immutable blocks* representing constructors and immutable records (as opposed to *mutable blocks* representing mutable records), e.g. `Some 0`. The physical comparison of `Some 0` and `Some 0` may return either `true` or `false`: we cannot determine the result of physical comparison just by looking at the abstract values.

To solve these problems we treat physical equality on abstract values is *non-deterministic* — even though the comparison instruction that implements it on low-level values is perfectly deterministic. The question is then: what guarantees do we get when physical equality returns `true` and when it returns `false`? Given such guarantees, denoted by `val_physeq` and `val_physneq`, the non-deterministic semantics is reflected in the logic through the following specification:

Lemma `physeq_spec v1 v2 :`

```

626 {{{ True }}}
627 v1 == v2
628 {{{ b, RET #b; 「(if b then val_physeq else val_physneq) v1 v2」 }}}

```

Proof. ... **Qed.**

The OCAML manual documents a partial specification for physical equality, which is precise for basic types such as integers or integer references, but does not clearly extend to structured values containing a mix of immutable and mutable constructors, which are present in the programs we verify. The only guarantee that it provides for all values is: if two values are physically equal, they are also structurally equal. For values that contain immutable constructors, we do not learn anything when they are physically distinct.

We will now explore the specifications of the `true` and `false` return cases. We describe our program verification requirements to suggest a precise enough semantics: if it does not say enough, we cannot prove our programs correct. In the other direction we describe some optimizations that OCAML implementations may perform (gathered through our discussion with OCAML maintainers), in particular the current compiler, that may rule out certain stronger specifications are incorrect.

Remark. It is extremely tempting to state that physical equality implies equality of `Rocq` representations, but incorrect in general as we have seen. Existing work on the modelling of OCAML physical equality within proof assistants have typically made this simplification, restricting the set of supported values to preserve soundness. We discussed our examples with authors of such earlier formalizations, and this has sometimes uncovered soundness issues on their side, as we discuss in Section 12.3.

5.3 When physical equality returns `true`

Let us go back to the concurrent stack of Figure 1 and more specifically the push function. Its atomic specification, given in Section 2.3, states that if we push a value v onto a stack whose current model is vs , then it atomically becomes a stack of model $v :: vs$. To prove this specification we rely on the fact that, if `Atomic.compare_and_set` returns `true`, the current list must be equal to vs , in the sense of `Rocq` equality for (abstract) Zoo values. This equality is strictly stronger than structural equality on mutable types, so we need a more precise specification than provided by the OCAML manual.

Zoo supports the following fragment of OCAML values: booleans, integers, mutable blocks (pointers), immutable blocks, functions.

The easy cases are mutable blocks and functions. Each of these two classes is disjoint from all others. We can reasonably assume that, when physical equality returns `true` and one of the compared values belongs to either of these classes, the two values are actually the same in `Rocq`. As far as we are aware, there is no optimization that could break this.

Booleans, integers and empty immutable blocks (constant constructors) are all represented as immediate integers in OCAML's low-level representation. This encoding induces conflicts: two abstract values that are distinct in `Rocq` may have the same low-level representation. The semantics of unrestricted physical equality has to reflect these conflicts: on those values that have an immediate representation, our specification does not state that physical equality includes equality of abstract values, it introduces a (simplified) notion of low-level representation and only states that those representations are equal.

Finally, let us consider the case of non-empty immutable blocks. At runtime, they are represented by pointers to tagged memory blocks. At first approximation, it is tempting to say that physically equal immutable blocks really are definitionally equal in `Rocq`. Alas, this is not true. To explain why, we have to recall that the OCAML compiler may perform *sharing*: immutable blocks containing physically equal fields may be shared. For example, the following tests may return `true`:

```
let test1 = Some 0 == Some 0 (* true *)
let test2 = [0;1] == [0;1] (* true *)
```

On its own, sharing is not a problem. However, coupled with representation conflicts, it can be surprising. Indeed, consider the any type we introduced previously:

```
type any = Any : 'a -> any
```

The following tests may return `true` (they do with `ocaml_opt`, not `ocamlc`):


```

687 let test1 = Any false == Any 0 (* true *)
688 let test2 = Any None == Any 0 (* true *)
689 let test3 = Any [] == Any 0 (* true *)

```

Now, going back to the push function of Figure 1, we have a problem. Given a stack of any, it is possible for the `Atomic.compare_and_set` to observe a current list (e.g. the one-element list `[Any 0]`) physically equal to the expected list (e.g., `[Any false]`) while these are actually distinct in Rocq. In short, the expected specification of Section 2.3 is incorrect: we may not get $v :: vs$ back in the model, but a list $v :: vs'$ where vs' is physically equal to vs but not the same abstract value. To fix this discrepancy, we would need to weaken all our specifications to be formulated *modulo physical equality*, which is non-standard and quite burdensome.

We believe this really is a shortcoming, at least from the verification perspective. Therefore, we propose to extend OCAML with *generative immutable blocks*⁸. These generative blocks are just like regular immutable blocks, except they cannot be shared. Hence, if physical equality on two generative blocks returns `true`, these blocks are definitionally equal in Rocq. At user level, this notion is materialized by *generative constructors*. For instance, to verify the expected push specification, we can use a generative version of lists:

```

703 type 'a glist =
704   | Nil
705   | Cons of 'a * 'a glist [@generative]

```

We modified the Zoo translator to support those generative constructors, and modified our implementation of Treiber stack to use the type `'a glist` instead of `'a list`, so that we could finally prove the expected, convenient specification. The `[@generative]` attribute is ignored by the standard OCAML compiler, and we have an experimental version that disables sharing for generative constructors.

5.4 When physical equality returns false

Most formalizations of physical equality in the literature do not give any guarantee when physical equality returns `false`. Many use-cases of physical equality, in particular retry loops, can be verified with only sufficient conditions on `true`. However, in some specific cases, more information is needed.

Consider the `Rcfd` module from the `Eio` library, an excerpt of which is given in Figure 3⁹. Thomas Leonard, its author, suggested that we verify this real-life example because of its intricate logical state. However, we found out that it is also relevant regarding the semantics of physical equality. Essentially, it consists in wrapping a file descriptor in a thread-safe way using reference-counting. At creation in the `make` function, the wrapper starts in the `Open` state. At some point, it can switch to the `Closing` state in the `close` function, and will remain `Closing` forever. Crucially, the `Open` state changes at most once to `Closing`, never to another `Open`.

The interest of `Rcfd` lies in the `close` function. First, the function reads the state. If this state is `Closing`, it returns `false`; the wrapper has been closed. If this state is `Open`, it tries to switch to the `Closing` state using `Atomic.Loc.compare_and_set`; if this attempt fails, it also returns `false`. In this particular case, we would like to prove that the wrapper has been closed, or equivalently that `Atomic.Loc.compare_and_set` cannot have observed `Open`. Intuitively, if we observed a different value then it must be `Closing`.

Obviously, we need some kind of guarantee related to the *physical identity* of `Open` when `Atomic.Loc.compare_and_set` returns `false`. If `Open` were a mutable block, we could argue that

⁸Non-anonymous link: <https://<clickable>>

⁹We make use of *atomic record fields* as introduced in Section 4.1.

```

736 type state =
737   | Open of Unix.file_descr
738   | Closing of (unit -> unit)
739
740 type t =
741   { mutable ops: int [@atomic];
742     mutable state: state [@atomic]; }
743
744 let make fd = { ops = 0; state = Open fd }
745
746 let closed = Closing (fun () -> ())
747 let close t =
748   match t.state with
749   | Closing _ -> false
750   | Open fd as prev ->
751     let next = Closing (fun () -> Unix.close fd) in
752     if Atomic.Loc.compare_and_set [%atomic.loc t.state] prev next then (
753       if t.ops == 0
754       && Atomic.Loc.compare_and_set [%atomic.loc t.state] next closed
755       then close () ;
756       true
757     ) else false
758
759
760
761
762
763
764
765
766
767
768
769
770
771
772
773
774
775
776
777
778
779
780
781

```

Fig. 3. Rcf module from Eio (excerpt)

this block cannot be physically distinct from itself; no optimization we know of would allow that. Unfortunately, it is an immutable block, and immutable blocks are subject to more optimizations. In fact, something surprising but allowed¹⁰ by OCAML can happen: *unsharing*, the dual of sharing. Indeed, any immutable block can be unshared, that is reallocated as if its definition was inlined. For example, the following test may theoretically return *false*:

```

767 let x = Some 0
768 let test = x == x (* false *)

```

Going back to *Rcf*, we have a problem: in the second branch, the *Open* block corresponding to *prev* could be unshared, which would make *Atomic.Loc.compare_and_set* fail. Hence, we cannot prove the expected specification; in fact, the program as it is written has a bug.

To remedy this unfortunate situation, we propose to reuse the notions of generative immutable blocks, that we introduced to prevent sharing, to also forbid unsharing by the OCAML compiler — we implemented this in our experiment compiler branch.

Supporting this requires enriching the Zoo semantics so that each generative block is annotated with a *logical identifier*¹¹ representing its physical identity, much like a pointer for a mutable block. If physical equality on two generative blocks returns *false*, the two identifiers are necessarily distinct. Given this semantics, we can verify the *close* function. Indeed, if *Atomic.Loc.compare_and_set* fails, we now know that the identifiers of the two blocks, if any, are distinct. As the concurrent protocol has only one *Open* block whose identifier does not change, it cannot be the case that the

¹⁰OCAML maintainers developing the *FLAMBDA* optimiser have confirmed that it may perform unsharing in certain cases.

¹¹Actually, for practical reasons, we distinguish identified and unidentified generative blocks.

current state is **Open**, hence it is **Closing**. We can verify this function after adding the following annotation:

```
type state =
  | Open of Unix.file_descr [@generative]
  | Closing of (unit -> unit)
```

5.5 Summary

In summary, we extended our abstract values with generative immutable blocks, and give the following specification to physical equality in ZOO_{LANG}. It can also serve as a precise specification of physical equality of a practical fragment of OCAML:

- On values whose low-level representation is an immediate integer, physical equality is immediate equality.
- On mutable blocks at some location, or generative immutable blocks with some identity, physical equality is equality of locations or identities.
- On immutable blocks, physical-equality is under-specified, but it implies that the blocks have the same tags and their arguments are in turn physically equal.
- Two values that fall into different categories above are never physically equal.

We reflect this specification in the Zoo program logic, while keeping a reasonably high-level definition of abstract values.

Our verification work uncovered correctness bugs in existing OCAML concurrent code due to unsharing. We propose to extend the language with a `[@generative]` annotation for immutable constructors whose physical identity we want to reason about when implementing a fine-grained concurrent data structure. A rule of thumb would be to use it for all non-constant constructors involved in **Atomic**.compare_and_set operations.

6 STRUCTURAL EQUALITY

Structural equality is also supported. More precisely, it is not part of the semantics of the language but implemented using low-level primitives¹². The reason is that it is in fact difficult to specify for arbitrary values. In general, we have to compare graphs — which implies structural comparison may diverge.

Accordingly, the specification of $v_1 = v_2$ requires the (partial) ownership of a *memory footprint* corresponding to the union of the two compared graphs, giving the permission to traverse them safely. If it terminates, the comparison decides whether the two graphs are bisimilar (modulo representation conflicts, as described in Section 5). In IRIS, this gives:

```
Lemma structeq_spec v1 v2 footprint,
  val_traversable footprint v1 →
  val_traversable footprint v2 →
  {{{ structeq_footprint footprint }}}
  v1 = v2
  {{{ b, RET #b;
    structeq_footprint footprint *
    ⌈(if b then val_structeq else val_structneq) footprint v1 v2⌋ }}}.
```

Obviously, this general specification is not very convenient to work with. Fortunately, for abstract values (without any mutable part), we can prove a much simpler variant saying that structural equality boils down to physical equality:

¹²In OCAML, these primitives correspond to the unsafe functions `Obj.is_int`, `Obj.tag`, `Obj.size` and `Obj.field`.

```

834 Lemma structeq_spec_abstract v1 v2 :
835   val_abstract v1 →
836   val_abstract v2 →
837   {{{ True }}}
838   v1 = v2
839   {{{ b, RET #b; ⌈(if b then val_physeq else val_physneq) v1 v2⌋ }}}
840 Proof. ... Qed.

```

Gabriel Je ne trouve pas du tout naturelle la dernière spécification. Lue naïvement elle suggère que (1) si deux valeurs immutables sont structurellement égales, alors elles sont physiquement égales, et que (2) si elles sont structurellement distinctes alors elles sont physiquement distinctes. La suggestion (1) est incorrecte et la suggestion (2) est trop faible. Ce qui se passe c'est que `val_physeq`, `neq` veulent dire autre chose que l'égalité physique, et sont des conséquences à la fois de l'(in)égalité physique et de l'(in)égalité structurelle. Je pense qu'il vaudrait mieux changer leurs noms.}

7 STANDARD DATA STRUCTURES

To save users from reinventing the wheel, we provide a library of verified standard data structures — more or less a subset of the OCAML standard library. Most of these data structures¹³ are completely reimplemented in Zoo and axiom-free, including the `Array`¹⁴ module.

Sequential data structures. We provide verified implementations of various sequential data structures: array, dynamic array (vector), list, stack, queue (bounded and unbounded), double-ended queue. We claim that the proven specifications are modular and practical. In fact, most of these data structures have already been used to verify more complex ones — we present some in [Section 8](#) and [Section 10](#). Especially, we developed an extensive collection of flexible specifications for the iterators of the `Array` and `List` modules. Remarkably, our formalization of `Array` features different (fractional) predicates to express the ownership of either an entire array, a slice or even a circular slice — we use it to verify algorithms involving circular arrays, e.g. Chase-Lev working-stealing queue [[Chase and Lev 2005](#)] as presented in [Section 10.4](#).

Concurrent data structures. We provide verified implementations of various concurrent data structures: domain¹⁵ (including domain-local storage), mutex, semaphore, condition variable, write-once variable (also known as *ivar*), atomic array. Note that there is currently no `Atomic_array` module in the OCAML standard library, but we are planning to propose it.

8 PERSISTENT DATA STRUCTURES

To further demonstrate the practicality of Zoo, we verified a collection of persistent data structures. This includes purely functional data structures such as persistent stack and queue, but also efficient imperative implementations of persistent array from [Conchon and Filliâtre \[2007\]](#), store and union-find from [Allain, Clément, Moine and Scherer \[2024\]](#).

Currently, verification of purely functional programs relies on the regular ZOOLANG translation, i.e. on a deeply embedded representation. However, we found this approach is cumbersome. In the future, it would be desirable to be able to verify them directly in [Rocq](#), through a translation to GALLINA. Similarly to HACSPEC [[Haselwarter, Hvass, Hansen, Winterhalter, Hritcu and Spitters](#)

¹³For practical reasons, to make them completely opaque, we chose to axiomatize a few functions from the `Domain` and `Random` modules. They could trivially be realized in Zoo.

¹⁴Our implementation of the `Array` module is compatible with the standard one. In particular, it uses the same low-level value representation.

¹⁵Domains are the units of parallelism in OCAML 5.

$$\begin{array}{c}
\text{883} \quad \text{inv } t \text{ fd } \Psi \\
\text{884} \quad \frac{\Psi \text{ 1}}{\text{make fd}} \\
\text{885} \quad \frac{}{t. \text{ inv } t \text{ fd } \Psi} \\
\text{886} \\
\text{887} \quad \frac{\text{wp closed } () \{ X \text{ false } \} \quad \forall q. \Psi \text{ } q \text{ } \neg * \text{ wp open fd } \{ \text{res}. \Psi \text{ } q * X \text{ true res } \}}{\text{use } t \text{ closed open}} \\
\text{888} \quad \frac{}{\text{res}. \exists b. X \text{ } b \text{ res}} \\
\text{889} \\
\text{890} \quad \frac{\text{inv } t \text{ fd } \Psi \quad \Psi \text{ 1 } \neg * \exists \text{ chars. unix-fd fd 1 chars}}{\text{close } t} \\
\text{891} \quad \frac{}{(). \text{ closing } t} \\
\text{892} \\
\text{893} \quad \frac{\text{inv } t \text{ fd } \Psi}{\text{remove } t} \\
\text{894} \quad \frac{}{0. \quad o = \text{None} \vee o = \text{Some fd } * \Psi \text{ 1}} \\
\text{895} \quad \text{closing } t
\end{array}$$

Fig. 4. Rcfd specification (excerpt)

2024], this new translation would come with a generated proof of equivalence with the ZooLang representation.

9 RCDF: PARALLELISM-SAFE FILE DESCRIPTOR

As mentioned in Section 5, the **Rcfd** module from the **Eio** library is particularly interesting in several respects. Not only does it justify the introduction of generative constructors in OCAML, but it also demonstrates the use of **Iris** for expressing realistic concurrent protocols.

Specification. The **Rcfd** module provides a parallelism-safe file descriptor (FD) relying internally on reference-counting. Interestingly, it is used in **Eio** in two different ways, more precisely two different ownership regimes: 1) any thread can try to access or close the FD; 2) any thread can try to access the FD but only the owner thread can close it — and is responsible for closing it. To verify all uses, the specification of **Rcfd** has to support both ownership regimes. However, due to space constraints, we consider a simplified specification given in Figure 4. The full verified specification can be found in the mechanization.

The specification features four operations¹⁶.

make creates a new object t of type **Rcfd**. t wrapping a given FD fd , yielding the (persistent) invariant $\text{inv } t \text{ fd } \Psi$, where Ψ is an arbitrary fractional predicate. Crucially, the user must provide the full predicate $\Psi \text{ 1}$, which is stored in the invariant. Once it is created, a wrapped FD can be accessed through the **use** operation and closed through the **close** operation.

use requires the invariant along with the weakest preconditions of the *closed* function, that is called if the FD has been flagged as closed, and *open* function, that is called if the FD is still open. To control the postconditions of the specification and the weakest preconditions, the user can choose an arbitrary predicate X parameterized by a boolean indicating whether the *closed* (**false**) or the *open* (**true**) was called. Furthermore, the *open* function is given a fraction of Ψ during its execution, thereby accessing the FD.

close requires the invariant and proving that the full ownership of Ψ entails the full ownership of the FD fd , which is necessary to call **Unix**. **close**. It yields **closing** t , a persistent resource witnessing that t has been flagged as closed. Actually, the wrapped FD is not closed immediately. It will be closed only once it is possible, meaning all ongoing calls to **use** owning a fraction of the FD end.

Alternatively, instead of closing the FD, **remove** tries to retrieve the full ownership of Ψ . To achieve it, it exploits the same mechanism as **close** — flagging t as closed as witnessed by **closing** t — but also waits until all use calls are done.

¹⁶We omitted two non-essential operations: **is_open** and **peek**.

Logical state. Thomas Leonard, the author of [Rcfd](#), suggested verifying it to make sure the informal concurrent protocol he described in the OCAML interface was correct. This protocol introduces a notion of monotonic logical state — modeled in [Iris](#) using a specific resource algebra [Timany and Birkedal 2021] — to describe the evolution of a FD. Originally, there were four logical states but we found that only three are necessary for the verification: **open**, **closing/users** and **closing/no-users**.

In the **open** state, the FD is available for use, meaning any thread can access it through use. Physically, this corresponds to the **Open** constructor.

When some thread flags the FD as closed through `close` or `remove`, the state transitions from **open** to **closing/users**. Crucially, there can only be one such thread. In this state, the FD is not really closed yet because of ongoing use operations. Physically, this logical transition corresponds to switching from the **Open** to the **Closing** constructor using `Atomic.Loc.compare_and_set`.

Once all use operations have finished, when the reference-count reaches zero, it is time to actually ‘close’ the FD by calling the function carried by the **Closing** constructor. However, this has to be done only once. The ‘closing’ thread is the one that succeeds in updating the **Closing** constructor (to a new one carrying a no-op function) using again `Atomic.Loc.compare_and_set`. At this point, the state transitions from **closing/users** to **closing/no-users** and the wrapper no longer owns the FD.

Generative constructors. As explained in [Section 5](#), the **Open** constructor has to be generative to prevent *unsharing*. In fact, the **Closing** constructor also has to be generative to prevent *sharing*: when the ‘closing’ thread wins the `Atomic.Loc.compare_and_set`, it must be true that the ‘closing’ function has not changed.

10 SATURN: A LIBRARY OF STANDARD LOCK-FREE DATA STRUCTURES

We verified a collection of standard lock-free data structures from the [Saturn](#), [Eio](#) and [Picos](#) [Karvonen 2025b] libraries. It includes stacks, queues (list-based, array-based and stack-based), bags, and work-stealing queues. These data structures are meant to be used as is or adapted to fit specific needs. To cover a wide range of use cases, we provide specialized variants: bounded or unbounded, single-producer (SP) or multi-producer (MP), single-consumer (SC) or multi-consumer (MC).

Due to space constraints, we focus on the most important algorithms and refrain from showing the corresponding (non-trivial) [Iris](#) invariants, which are mechanized in [Rocq](#).

10.1 Stacks

We verified three variants of the Treiber stack [Treiber 1986]: 1) unbounded MPMC (the standard one), 2) bounded MPMC, 3) closable unbounded MPMC. This last variant features a closing mechanism: at some point, some thread can decide to close the stack, retrieving the current content and preventing others from operating on it. For example, we used it to represent a set of vertex successors in the context of a concurrent graph implementation (not presented in this paper).

As explained in [Section 5](#), the three verified stacks use generative constructors to prevent sharing. One may ask whether it would be easier to use a mutable version of lists instead. From the programmer’s perspective, this is unsatisfactory because 1) the compiler will typically emit warnings complaining that the mutability is not exploited and 2) it does not really reflect the intent, *i.e.* we want precise guarantees for physical equality, not modify the list. From the verification perspective, this is also unsatisfactory because the mutable representation is more complex to write and reason about: pointers and points-to assertions versus pure [Rocq](#) list.

Although verified stacks may seem like a not-so-new contribution, it is, as far as we know, the first verification of realistic OCAML implementations. For comparison, the exemplary concurrent

stacks verified in [IRIS](#) [[Iris development team 2025](#)] all suffer from the same flaw: they need to introduce indirections (pointers) to be able to use the compare-and-set primitive.

10.2 List-based queues

We verified four variants of the Michael-Scott queue [[Michael and Scott 1996](#)]: unbounded MPMC (the standard one), bounded MPMC, unbounded MPSC and unbounded SPMC. The SPMC queue is used to implement one of the bags, more precisely a relaxed queue guaranteeing only per-producer ordering.

In the [IRIS](#) literature, [Vindum and Birkedal \[2021\]](#) established contextual refinement of the Michael-Scott queue while [Mulder and Krebbers \[2023\]](#) proved logical atomicity. However, we had to redesign and extend the invariant for several reasons.

Efficient implementation. The Michael-Scott essentially consists of a singly linked list of nodes that only grows over time. The previously verified implementations, implemented in [HEAPLANG](#), use a double indirection to represent the list [[Vindum and Birkedal 2021](#), Figure 2]. Similarly to the Treiber stack, this is made so as to be able to use the compare-and-set primitive of [HEAPLANG](#). [Vindum and Birkedal \[2021\]](#) write:

A node is a pointer to either none or some of a pair of a value and a pointer to the next node. The pointer serves to make nodes comparable by pointer equality such that pointers to nodes can be changed with CAS.

In OCAML, this would correspond to introducing extra atomic references ([Atomic.t](#)) between the nodes. Using atomic record fields, we can represent the list more efficiently, without the extra indirection. However, there is one subtlety: in this new representation, we need to clear the outdated nodes so that their value is no longer reachable and can be garbage-collected, in other words to prevent a memory leak. This subtlety was missing from the original implementation designed for non-garbage-collected languages, but [Saturn](#) implements it ([saturn#64](#)).

To deal with this representation in separation logic, we introduce the notion of *explicit chain* that allows decoupling the chain structure formed by the nodes and the content of the nodes. Concretely, the assertion $xchain\ dq\ \ell s\ dst$ represents a chain linking locations ℓs and ending at value dst ; dq is a discardable fraction [[Vindum and Birkedal 2021](#)] that controls the ownership of the chain. This notion is very flexible as it is independent of the rest of the structure. As a matter of fact, we used it and its generalization to doubly linked list more broadly, to verify other algorithms. All the variants of Michael-Scott we verified rely on it. In particular, it was quite straightforward to extend the invariant of the bounded queue, where nodes carry more (mutable and immutable) information.

External linearization point. Our work also revealed another interesting aspect that is not addressed in the literature, as far as we know. None of the previously verified implementations deal with the `is_empty` operation, that consists in reading the sentinel node and checking whether it has a successor. If it has no successor, it is necessarily the last node of the chain, hence the queue is empty. If it does have a successor, `is_empty` returns `false`, meaning we must have observed a non-empty queue. However, this last part is more tricky than it may seem. Indeed, it may happen that 1) we read the sentinel while the queue is empty, 2) other operations fill and empty again the queue so that the sentinel is outdated, 3) we read the successor of the former sentinel while the queue is still empty. The crucial point here is that `is_empty` is linearized when the first push operation filled the queue. In other words, the linearization point of `is_empty` is triggered by another operation; this is called an *external linearization point*. To handle this in the proof, we

introduced a mechanism in the invariant to transfer the [Iris](#) resource materializing the linearization point¹⁷ from `is_empty` to `push` and vice versa.

10.3 Stack-based queues

10.4 Work-stealing queues

11 MEMORY SAFETY

Concurrency creates new sources of tension between performance and memory-safety. The OCAML maintainers are trying to maintain memory-safety for all OCAML programs, including concurrent programs that would contain data races. They have tried to imprint this focus on memory-safety to library authors as well.

Performance-sensitive OCAML libraries are sometimes written using unsafe primitives, that may break memory-safety if used incorrectly. It is the responsibility of code authors to ensure that the preconditions of those unsafe primitives are satisfied to ensure safety. Unfortunately, the addition to parallel code execution in OCAML 5 broke the safety of some existing unsafe code: it adds more possible interleaving and may invalidate safety reasoning.

11.1 Dynarray

In January 2023, Gabriel Scherer proposed the additional of [Dynarray](#), a module of (sequential) resizable arrays, to the OCAML standard library.¹⁸

A concurrent safety problem. Resizable arrays are implemented as a record with two mutable fields, a size field that stores the current length of the resizable array, and a data field storing the “backing array” a (non-resizable) array of size at least `size` elements. When a user pushes a new element to the end of a resizable array, it typically suffices to write the new element at index `size` in the backing array, and then increment the size field. But when the size field reaches the actual length of the backing array (which we call the “capacity” of the resizable array), we need to allocate a new, larger backing array, to copy the values from the old to the new backing array, and to overwrite the data field with the new backing array; then we can write the user-provided element.

This operation of pushing a new element is performance-critical for resizable arrays, and it always incurs a size comparison followed by resizing is necessary. After this check and potential resizing, sequentially we know that the backing array has enough space, and we can write the new element using an `unsafe_set` operation that avoids a second bound-checking on array access. Efficient implementations use this unsafe operation, which can provide performance benefits of up to 20% in certain scenarios. Unfortunately, the safety reasoning is incorrect in presence of parallel operations on the same dynarray: another thread could mutate the backing array after the size check and before the unsafe write, for example [Dynarray](#).`reset` which sets the size to 0 and replaces the backing array by an empty array.

```
let add_last a x = (* sequential version *)
  let size = a.size in
  if Array.length a.data = size then ensure_capacity a (size + 1);
  a.size <- size + 1;
  Array.unsafe_set a.arr size x
```

Writing to a (non-resizable) array outside its bounds without bound checks breaks memory safety, so this implementation was memory-safe under OCAML 4 but it becomes memory-unsafe under

¹⁷This resource is known as an *atomic update*. [Mulder and Krebbers \[2023\]](#) provide a good description.

¹⁸<https://github.com/ocaml/ocaml/pull/11882>

OCAML 5. On the other hand, **Dynarray** is explicitly documented as a sequential data structure, so it is the user responsibility to respect this precondition by using appropriate synchronization (for example a mutex) to prevent data races. Some performance-obsessed users would argue that if *other* users fail in their responsibility of ensuring sequential access, then they do not deserve memory-safety. The OCAML maintainers and standard library authors consider on the other hand that memory-safety should be preserved even in this case: it can only be lifted for operations that are explicitly marked as unsafe, and hopefully have simple, checkable preconditions. So they decided to change the optimization of **Dynarray** to guarantee memory-safety even in presence of racy concurrent usage.

The proposed implementation reads the data field to get the backing array *b*, then it performs the size check on this backing array. If no resizing is necessary, then an unsafe write is performed as before (on the backing array *b*, without re-reading the data field). In the infrequent case where resizing is necessary, the operation retries afterwards. (A full retry is not necessary, a bound-checking write would suffice for safety.)

```
let rec add_last a x = (* memory-safe under concurrency *)
  let {size; data} = a in
  if Array.length data >= size
  then ( ensure_capacity a (size + 1); add_last a x )
  else ( a.size <- size + 1; Array.unsafe_set arr size x )
```

A concurrent verification problem. We verified (a representative fragment of) the proposed **Dynarray** implementation in ZOO_{LANG}, providing a definitive proof that it is functionally correct under sequential usage, and that it does preserve memory-safety even under concurrent usage. The verified fragment is now part of the Zoo standard library, and can be used in further verification projects.

Informally, the verification of (sequential) functional correctness and (concurrent) memory-safety rely on two different invariants:

- Functional correctness relies on a *strong* invariant, which may not be preserved under concurrent usage; for example that the content of the size field is smaller than the size of the backing array.
- Memory-safety relies on a *weak* invariant that does not suffice to prove correctness, but is preserved under concurrent usage. For this **Dynarray** implementation, the weak invariant is that the backing array remains well-typed and that the content of the size field is non-negative.

To informally check the two desired properties, it suffices to check that the strong invariant is indeed a sequential invariant, and that implies functional correctness; and separately that the weak invariant is preserved by all operations on the data structure.

In our mechanized verification, the strong invariant becomes a *model* of the data structure, and we verify separation-logic triples where each operation can assume unique ownership of the model in input, and has to return ownership of a valid model in output, as is standard. For example, the “functional correctness” lemma for the `add_last` function is the following:

```
Lemma dynarray_2_push_spec t vs v :
  {{{ dynarray_2_model t vs }}}
  dynarray_2_push t v
  {{{ RET (); dynarray_2_model t (vs ++ [v]) }}}.
```

This lemma establishes correctness in a purely-sequential usage, but also under any concurrent usage where correctly synchronization is used to guarantee unique ownership of the model.

The formal counterpart of our weak invariant is a *semantic type*, following the approach of the RUSTBELT project [Jung, Jourdan, Krebbers and Dreyer 2018], which also deals with unsafe fragments within a language intended to be safe. Note that the RUSTBELT semantic types are derived from RUST types, where mutating operations get a mutable borrow and thus unique ownership (for a time). In contrast, our semantic types carry no ownership of the values they manipulate, so they are stored entirely as invariants, and all interactions with the structure must be shown to preserve this invariant atomically. The invariant must be robust against any interference coming from any other function called in parallel on the same structure.

The definition of the semantic type for dynarrays, and the statement of memory safety (or in fact semantic typing) for `add_last` look as follows:

```

Definition itype_dynarray_2 ty t : iProp  $\Sigma$  :=
   $\exists$  l,
   $\ulcorner t = \#l \urcorner *$ 
  inv nroot (
     $\exists$  (sz : nat) cap data,
    l.[size] #sz *
    l.[data] data * itype_array ty cap data
  ).

```

```

Lemma dynarray_2_push_type ty t v :
  {{{ itype_dynarray_2 ty t * ty v }}}
  dynarray_2_push ty t v
  {{{ RET (); True }}}.

```

The predicate `ty` is the semantic type for the elements of the array – it must hold for each value of the backing array and for the new value `v`.

Gabriel{Je crois avoir compris qu'il y a des choses intéressantes à dire sur le raisonnement sur les itérateurs (`iteri` par exemple), mais je ne saurais pas quoi écrire à ce sujet.}

11.2 Saturn: Single-producer or Single-consumer queues

12 RELATED WORK

In general there are two approaches to practical program verification:

12.1 Non-automated verification

The verified program is translated, manually or in an automated way, into a representation living inside a proof assistant. The user has to write specifications and prove them.

The representation may be primitive, like Gallina for Rocq. For pure programs, **Gabriel**[this is rather straightforward]{I disagree, I believe that hs-to-coq is scientifically problematic as the translation is unsound for higher-order functions and infinite data.}, e.g. in hs-to-coq [Spector-Zabusky, Breitner, Rizkallah and Weirich 2018]. For imperative programs, this is more challenging. One solution is to use a monad, e.g. in coq-of-ocaml [Claret 2025], but it does not support concurrency.

The representation may be embedded, meaning the semantics of the language is formalized in the proof assistant. This is the path taken by some recent works [Chajed, Tassarotti, Kaashoek and Zeldovich 2019; Charguéraud 2023; Daby-Seesaram, Madiot, Pottier, Seassau and Yoon 2024; Gondelman, Hinrichsen, Pereira, Timany and Birkedal 2023] harnessing the power of separation logic. In particular, CFML [Charguéraud 2023] and OSIRIS [Daby-Seesaram, Madiot, Pottier, Seassau

and Yoon 2024] target OCAML. However, CFML does not support concurrency and is not based on IRIS. OSIRIS, still under development, is based on IRIS but does not support concurrency.

At the time of writing, HEAPLANG is thus the most appropriate tool to verify concurrent OCAML programs. We discussed limitations of HEAPLANG in the introduction, and ZOO LANG is our proposal to improve on this. Conversely, one notable limitation of ZOO LANG today is its lack of support for OCAML’s relaxed memory model.

12.2 Semi-automated verification

In semi-automated verification approaches, the verified program is annotated by the user to guide the verification tool: preconditions, postconditions, invariants, *etc.* Given this input, the verification tool generates proof obligations that are mostly automatically discharged. One may further distinguish two types of semi-automated systems: *foundational* and *non-foundational*.

In *non-foundational* automated verification, the tool and the external solvers it may rely on are part of the trusted computing base. It is the most common approach and has been widely applied in the literature [Astrauskas, Bilý, Fiala, Grannan, Matheja, Müller, Poli and Summers 2022; Denis, Jourdan and Marché 2022; Filliâtre and Paskevich 2013; Jacobs, Smans, Philippaerts, Vogels, Penninckx and Piessens 2011; Lattuada, Hance, Cho, Brun, Subasinghe, Zhou, Howell, Parno and Hawblitzel 2023; Müller, Schwerhoff and Summers 2017; Pulte, Makwana, Sewell, Memarian, Sewell and Krishnaswami 2023; Swamy, Chen, Fournet, Strub, Bhargavan and Yang 2013], including to OCAML by CAMELEER [Pereira and Ravara 2021], which uses the GOSPEL specification language [Charguéraud, Filliâtre, Lourenço and Pereira 2019] and WHY3 [Filliâtre and Paskevich 2013].

In *foundational* automated verification, the proofs are checked by a proof assistant like Rocq, meaning the automation does not have to be trusted. To our knowledge, it has been applied to C [Sammler, Lepigre, Krebbers, Memarian, Dreyer and Garg 2021] and RUST [Gäher, Sammler, Jung, Krebbers and Dreyer 2024].

Zoo is a non-automated verification framework — except for our use DIAFRAME for local automation of separation logic reasoning. We would be interested in moving towards more automation in the future.

12.3 Physical equality

There is some literature in proof-assistant research on reflecting physical equality from the implementation language into the proof assistant, for optimization purposes: for example, exposing OCAML’s physical equality as a predicate in Rocq lets us implement some memoization and sharing techniques in Rocq libraries. However, axiomatizing physical equality in the proof assistant is difficult, and can result in inconsistencies.

The earlier discussions of this question that we know come from Jourdan’s thesis [Jourdan 2016] (chapter 9), also presented more succinctly in [Braubant, Jourdan and Monniaux 2014]. This work introduces the Jourdan condition, that physical equality implies equality of values. [Boulmé 2021] extends the treatment of physical equality in Rocq, integrating it in an “extraction monad” to control it more safely. There is also a discussion of similar optimizations in LEAN in [Selsam, Hudon and de Moura 2020].

The correctness of the axiomatization of physical equality depends on the type of the values being compared: axiomatizations are typically polymorphic on any type A, but their correctness depends on the specific A being considered. For example, it is easy to correctly characterize physical on natural numbers, and other non-dependent types arising in Rocq verification projects. One difficulty in HEAPLANG and ZOO LANG is that they are untyped languages, their representation of \emptyset and `false` has the same type. But our remark that structural equality (in OCAML) does not

necessarily coincide with definitional equality (in [Rocq](#)) also applies to other [Rocq](#) types: our examples with an existential [Any](#) constructor (see [Section 5](#)) can be reproduced with Σ -types.

13 CONCLUSION AND FUTURE WORK

We presented Zoo, a framework for the verification of concurrent OCAML 5 programs. While it is not yet available on opam, it can be installed and used in other [Rocq](#) projects. We provide a minimal example¹⁹ demonstrating its use.

Zoo has already been used to verify sequential imperative algorithms [anonymous] and is currently being used to verify a library of lock-free data structures. Its main weakness so far is its memory model, which is sequentially consistent as opposed to the relaxed OCAML 5 memory model. It also lacks exceptions and algebraic effects, that we plan to introduce in the future.

Another interesting direction would be to combine Zoo with semi-automated techniques. Similarly to [WHY3](#), the simple parts of the verification effort would be done in a semi-automated way, while the most difficult parts would be conducted in [Rocq](#).

REFERENCES

- Clément Allain, Basile Clément, Alexandre Moine, and Gabriel Scherer. 2024. Snapshottable Stores. *Proc. ACM Program. Lang.* 8, ICFP (2024), 338–369. doi:10.1145/3674637
- Vytautas Astrauskas, Aurel Bilý, Jonás Fiala, Zachary Grannan, Christoph Matheja, Peter Müller, Federico Poli, and Alexander J. Summers. 2022. The Prusti Project: Formal Verification for Rust. In *NASA Formal Methods - 14th International Symposium, NFM 2022, Pasadena, CA, USA, May 24-27, 2022, Proceedings (Lecture Notes in Computer Science, Vol. 13260)*, Jyotirmoy V. Deshmukh, Klaus Havelund, and Ivan Perez (Eds.). Springer, 88–108. doi:10.1007/978-3-031-06773-0_5
- Lars Birkedal, Thomas Dinsdale-Young, Armaël Guéneau, Guilhem Jaber, Kasper Svendsen, and Nikos Tzevelekos. 2021. Theorems for free from separation logic specifications. *Proc. ACM Program. Lang.* 5, ICFP (2021), 1–29. doi:10.1145/3473586
- Sylvain Boulmé. 2021. *Formally Verified Defensive Programming (efficient Coq-verified computations from untrusted ML oracles)*. Accreditation to supervise research. Université Grenoble-Alpes. <https://hal.science/tel-03356701> See also <http://www-verimag.imag.fr/~boulme/hdr.html>.
- Thomas Braibant, Jacques-Henri Jourdan, and David Monniaux. 2014. Implementing and Reasoning About Hash-consed Data Structures in Coq. *J. Autom. Reason.* 53, 3 (2014), 271–304. doi:10.1007/S10817-014-9306-0
- Tej Chajed, Joseph Tassarotti, M. Frans Kaashoek, and Nickolai Zeldovich. 2019. Verifying concurrent, crash-safe systems with Perennial. In *Proceedings of the 27th ACM Symposium on Operating Systems Principles, SOSP 2019, Huntsville, ON, Canada, October 27-30, 2019*, Tim Brecht and Carey Williamson (Eds.). ACM, 243–258. doi:10.1145/3341301.3359632
- Yun-Sheng Chang, Ralf Jung, Upamanyu Sharma, Joseph Tassarotti, M. Frans Kaashoek, and Nickolai Zeldovich. 2023. Verifying vMVCC, a high-performance transaction library using multi-version concurrency control. In *17th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2023, Boston, MA, USA, July 10-12, 2023*, Roxana Geambasu and Ed Nightingale (Eds.). USENIX Association, 871–886. <https://www.usenix.org/conference/osdi23/presentation/chang>
- Arthur Charguéraud. 2023. *Habilitation thesis: A Modern Eye on Separation Logic for Sequential Programs. (Un nouveau regard sur la Logique de Séparation pour les programmes séquentiels)*. Université de Strasbourg. <https://tel.archives-ouvertes.fr/tel-04076725>
- Arthur Charguéraud, Jean-Christophe Filliâtre, Cláudio Lourenço, and Mário Pereira. 2019. GOSPEL - Providing OCaml with a Formal Specification Language. In *Formal Methods - The Next 30 Years - Third World Congress, FM 2019, Porto, Portugal, October 7-11, 2019, Proceedings (Lecture Notes in Computer Science, Vol. 11800)*, Maurice H. ter Beek, Annabelle McIver, and José N. Oliveira (Eds.). Springer, 484–501. doi:10.1007/978-3-030-30942-8_29
- David Chase and Yossi Lev. 2005. Dynamic circular work-stealing deque. In *SPAA 2005: Proceedings of the 17th Annual ACM Symposium on Parallelism in Algorithms and Architectures, July 18-20, 2005, Las Vegas, Nevada, USA*, Phillip B. Gibbons and Paul G. Spirakis (Eds.). ACM, 21–28. doi:10.1145/1073970.1073974
- Guillaume Claret. 2025. coq-of-ocaml. <https://github.com/formal-land/coq-of-ocaml>
- Sylvain Conchon and Jean-Christophe Filliâtre. 2007. A persistent union-find data structure. In *Proceedings of the ACM Workshop on ML, 2007, Freiburg, Germany, October 5, 2007*, Claudio V. Russo and Derek Dreyer (Eds.). ACM, 37–46. doi:10.1145/1292535.1292541

¹⁹Non-anonymous link: [https://<clickable>](#)

- Pedro da Rocha Pinto, Thomas Dinsdale-Young, and Philippa Gardner. 2014. TaDA: A Logic for Time and Data Abstraction. In *ECOOP 2014 - Object-Oriented Programming - 28th European Conference, Uppsala, Sweden, July 28 - August 1, 2014. Proceedings (Lecture Notes in Computer Science, Vol. 8586)*, Richard E. Jones (Ed.). Springer, 207–231. doi:10.1007/978-3-662-44202-9_9
- Arnaud Daby-Seesaram, Jean-Marie Madiot, François Pottier, Remy Seassau, and Irene Yoon. 2024. Osiris. <https://gitlab.inria.fr/fpottier/osiris>
- Xavier Denis, Jacques-Henri Jourdan, and Claude Marché. 2022. Creusot: A Foundry for the Deductive Verification of Rust Programs. In *Formal Methods and Software Engineering - 23rd International Conference on Formal Engineering Methods, ICFEM 2022, Madrid, Spain, October 24-27, 2022, Proceedings (Lecture Notes in Computer Science, Vol. 13478)*, Adrián Riesco and Min Zhang (Eds.). Springer, 90–105. doi:10.1007/978-3-031-17244-1_6
- Stephen Dolan, KC Sivaramakrishnan, and Anil Madhavapeddy. 2018. Bounding data races in space and time. *SIGPLAN Not.* 53, 4 (June 2018), 242–255. doi:10.1145/3296979.3192421
- Jean-Christophe Filliâtre and Andrei Paskevich. 2013. Why3 - Where Programs Meet Provers. In *Programming Languages and Systems - 22nd European Symposium on Programming, ESOP 2013, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2013, Rome, Italy, March 16-24, 2013. Proceedings (Lecture Notes in Computer Science, Vol. 7792)*, Matthias Felleisen and Philippa Gardner (Eds.). Springer, 125–128. doi:10.1007/978-3-642-37036-6_8
- Lennard Gäher, Michael Sammler, Ralf Jung, Robbert Krebbers, and Derek Dreyer. 2024. RefinedRust: A Type System for High-Assurance Verification of Rust Programs. *Proc. ACM Program. Lang.* 8, PLDI (2024), 1115–1139. doi:10.1145/3656422
- Léon Gondelman, Jonas Kastberg Hinrichsen, Mário Pereira, Amin Timany, and Lars Birkedal. 2023. Verifying Reliable Network Components in a Distributed Separation Logic with Dependent Separation Protocols. *Proc. ACM Program. Lang.* 7, ICFP (2023), 847–877. doi:10.1145/3607859
- Philipp G. Haselwarter, Benjamin Salling Hvass, Lasse Letager Hansen, Théo Winterhalter, Catalin Hritcu, and Bas Spitters. 2024. The Last Yard: Foundational End-to-End Verification of High-Speed Cryptography. In *Proceedings of the 13th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2024, London, UK, January 15-16, 2024*, Amin Timany, Dmitriy Traytel, Brigitte Pientka, and Sandrine Blazy (Eds.). ACM, 30–44. doi:10.1145/3636501.3636961
- Maurice Herlihy and Jeannette M. Wing. 1990. Linearizability: A Correctness Condition for Concurrent Objects. *ACM Trans. Program. Lang. Syst.* 12, 3 (1990), 463–492. doi:10.1145/78969.78972
- Iris development team. 2025. Iris examples. <https://gitlab.mpi-sws.org/iris/examples/>
- Bart Jacobs, Jan Smans, Pieter Philippaerts, Frédéric Vogels, Willem Penninckx, and Frank Piessens. 2011. VeriFast: A Powerful, Sound, Predictable, Fast Verifier for C and Java. In *NASA Formal Methods - Third International Symposium, NFM 2011, Pasadena, CA, USA, April 18-20, 2011. Proceedings (Lecture Notes in Computer Science, Vol. 6617)*, Mihaela Gheorghiu Bobaru, Klaus Havelund, Gerard J. Holzmann, and Rajeev Joshi (Eds.). Springer, 41–55. doi:10.1007/978-3-642-20398-5_4
- Jacques-Henri Jourdan. 2016. *Verasco: a Formally Verified C Static Analyzer. (Verasco: un analyseur statique pour C formellement vérifié)*. Ph. D. Dissertation. Paris Diderot University, France. <https://tel.archives-ouvertes.fr/tel-01327023>
- Ralf Jung, Jacques-Henri Jourdan, Robbert Krebbers, and Derek Dreyer. 2018. RustBelt: securing the foundations of the rust programming language. *PACMPL* 2, POPL (2018). doi:10.1145/3158154
- Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, Ales Bizjak, Lars Birkedal, and Derek Dreyer. 2018. Iris from the ground up: A modular foundation for higher-order concurrent separation logic. *J. Funct. Program.* 28 (2018), e20. doi:10.1017/S0956796818000151
- Ralf Jung, Rodolphe Lepigre, Gaurav Parthasarathy, Marianna Rapoport, Amin Timany, Derek Dreyer, and Bart Jacobs. 2020. The future is ours: prophecy variables in separation logic. *Proc. ACM Program. Lang.* 4, POPL (2020), 45:1–45:32. doi:10.1145/3371113
- Vesa Karvonen. 2025a. Kcas. <https://github.com/ocaml-multicore/kcas>
- Vesa Karvonen. 2025b. Picos. <https://github.com/ocaml-multicore/picos>
- Vesa Karvonen and Carine Morel. 2025. Saturn. <https://github.com/ocaml-multicore/saturn>
- Robbert Krebbers, Jacques-Henri Jourdan, Ralf Jung, Joseph Tassarotti, Jan-Oliver Kaiser, Amin Timany, Arthur Charguéraud, and Derek Dreyer. 2018. MoSeL: a general, extensible modal framework for interactive proofs in separation logic. *Proc. ACM Program. Lang.* 2, ICFP (2018), 77:1–77:30. doi:10.1145/3236772
- Andrea Lattuada, Travis Hance, Chanhee Cho, Matthias Brun, Isitha Subasinghe, Yi Zhou, Jon Howell, Bryan Parno, and Chris Hawblitzel. 2023. Verus: Verifying Rust Programs using Linear Ghost Types. *Proc. ACM Program. Lang.* 7, OOPSLA1 (2023), 286–315. doi:10.1145/3586037
- Anil Madhavapeddy and Thomas Leonard. 2025. Eio. <https://github.com/ocaml-multicore/eio>
- Glen Mével, Jacques-Henri Jourdan, and François Pottier. 2020. Cosmo: a concurrent separation logic for multicore OCaml. *Proc. ACM Program. Lang.* 4, ICFP (2020), 96:1–96:29. doi:10.1145/3408978
- Maged M. Michael and Michael L. Scott. 1996. Simple, Fast, and Practical Non-Blocking and Blocking Concurrent Queue Algorithms. In *Proceedings of the Fifteenth Annual ACM Symposium on Principles of Distributed Computing, Philadelphia, Pennsylvania, USA, May 23-26, 1996*, James E. Burns and Yoram Moses (Eds.). ACM, 267–275. doi:10.1145/248052.248106

- Ike Mulder and Robbert Krebbers. 2023. Proof Automation for Linearizability in Separation Logic. *Proc. ACM Program. Lang.* 7, OOPSLA1 (2023), 462–491. doi:10.1145/3586043
- Ike Mulder, Robbert Krebbers, and Herman Geuvers. 2022. Diaframe: automated verification of fine-grained concurrent programs in Iris. In *PLDI '22: 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation, San Diego, CA, USA, June 13 - 17, 2022*, Ranjit Jhala and Isil Dillig (Eds.). ACM, 809–824. doi:10.1145/3519939.3523432
- Peter Müller, Malte Schwerhoff, and Alexander J. Summers. 2017. Viper: A Verification Infrastructure for Permission-Based Reasoning. In *Dependable Software Systems Engineering*, Alexander Pretschner, Doron Peled, and Thomas Hutzelmann (Eds.). NATO Science for Peace and Security Series - D: Information and Communication Security, Vol. 50. IOS Press, 104–125. doi:10.3233/978-1-61499-810-5-104
- Mário Pereira and António Ravara. 2021. Cameleer: A Deductive Verification Tool for OCaml. In *Computer Aided Verification - 33rd International Conference, CAV 2021, Virtual Event, July 20-23, 2021, Proceedings, Part II (Lecture Notes in Computer Science, Vol. 12760)*, Alexandra Silva and K. Rustan M. Leino (Eds.). Springer, 677–689. doi:10.1007/978-3-030-81688-9_31
- Christopher Pulte, Dhruv C. Makwana, Thomas Sewell, Kayvan Memarian, Peter Sewell, and Neel Krishnaswami. 2023. CN: Verifying Systems C Code with Separation-Logic Refinement Types. *Proc. ACM Program. Lang.* 7, POPL (2023), 1–32. doi:10.1145/3571194
- Michael Sammler, Rodolphe Lepigre, Robbert Krebbers, Kayvan Memarian, Derek Dreyer, and Deepak Garg. 2021. RefinedC: automating the foundational verification of C code with refined ownership types. In *PLDI '21: 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation, Virtual Event, Canada, June 20-25, 2021*, Stephen N. Freund and Eran Yahav (Eds.). ACM, 158–174. doi:10.1145/3453483.3454036
- Daniel Selsam, Simon Hudon, and Leonardo de Moura. 2020. Sealing pointer-based optimizations behind pure functions. *Proc. ACM Program. Lang.* 4, ICFP, Article 115 (Aug. 2020), 20 pages. doi:10.1145/3408997
- KC Sivaramakrishnan, Stephen Dolan, Leo White, Sadiq Jaffer, Tom Kelly, Anmol Sahoo, Sudha Parimala, Atul Dhiman, and Anil Madhavapeddy. 2020. Retrofitting parallelism onto OCaml. *Proc. ACM Program. Lang.* 4, ICFP, Article 113 (Aug. 2020), 30 pages. doi:10.1145/3408995
- Antal Spector-Zabusky, Joachim Breitner, Christine Rizkallah, and Stephanie Weirich. 2018. Total Haskell is reasonable Coq. In *Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los Angeles, CA, USA, January 8-9, 2018*, June Andronick and Amy P. Felty (Eds.). ACM, 14–27. doi:10.1145/3167092
- Nikhil Swamy, Juan Chen, Cédric Fournet, Pierre-Yves Strub, Karthikeyan Bhargavan, and Jean Yang. 2013. Secure distributed programming with value-dependent types. *J. Funct. Program.* 23, 4 (2013), 402–451. doi:10.1017/S0956796813000142
- Amin Timany and Lars Birkedal. 2021. Reasoning about monotonicity in separation logic. In *CPP '21: 10th ACM SIGPLAN International Conference on Certified Programs and Proofs, Virtual Event, Denmark, January 17-19, 2021*, Catalin Hritcu and Andrei Popescu (Eds.). ACM, 91–104. doi:10.1145/3437992.3439931
- R. K. Treiber. 1986. *Systems Programming: Coping with Parallelism*. International Business Machines Incorporated, Thomas J. Watson Research Center. <https://books.google.fr/books?id=YQg3HAAACAAJ>
- Simon Friis Vindum and Lars Birkedal. 2021. Contextual refinement of the Michael-Scott queue (proof pearl). In *CPP '21: 10th ACM SIGPLAN International Conference on Certified Programs and Proofs, Virtual Event, Denmark, January 17-19, 2021*, Catalin Hritcu and Andrei Popescu (Eds.). ACM, 76–90. doi:10.1145/3437992.3439930
- Simon Friis Vindum, Dan Frumin, and Lars Birkedal. 2022. Mechanized verification of a fine-grained concurrent queue from meta's folly library. In *CPP '22: 11th ACM SIGPLAN International Conference on Certified Programs and Proofs, Philadelphia, PA, USA, January 17 - 18, 2022*, Andrei Popescu and Steve Zdancewic (Eds.). ACM, 100–115. doi:10.1145/3497775.3503689