Zoo: A framework for the verification of concurrent OCaml 5 programs using separation logic

ANONYMOUS AUTHOR(S)

The release of OCaml 5, which introduced parallelism in the OCaml runtime, drove the need for safe and efficient concurrent data structures. New libraries like `Saturn` address this need. This is an opportunity to apply and further state-of-the-art program verification techniques.

We present Zoo, a framework for verifying fine-grained concurrent OCaml 5 algorithms. Following a pragmatic approach, we defined a limited but sufficient fragment of the language to faithfully express these algorithms: ZooLang. We formalized its semantics carefully via a deep embedding in the Rocq proof assistant, uncovering subtle aspects of physical equality. We provide a tool to translate source OCaml programs into ZooLang syntax embedded inside Rocq, where they can be specified and verified using the Iris concurrent separation logic. To illustrate the applicability of Zoo, we verified a subset of the standard library and a collection of fined-grained concurrent data structures from the `Saturn` and `Eio` libraries.

In the process, we also extended OCaml to more efficiently express certain concurrent programs.

1 Introduction

OCaml 5.0 was released on December 15th 2022, the first version of the OCaml programming language to support parallel execution of OCaml threads by merging the Multicore OCaml runtime [Sivaramakrishnan, Dolan, White, Jaffer, Kelly, Sahoo, Parimala, Dhiman and Madhavapeddy 2020]. It provided basic support in the language runtime to start and stop coarse-grained threads called "domains", and support for strongly sequential atomic references in the standard library. The third-party library `domainslib` offered a simple scheduler for a pool of tasks, used to benchmark the parallel runtime. A world of parallel and concurrent software was waiting to be invented.

Shared-memory concurrency is a difficult programming domain, and existing ecosystems (C++, Java, Haskell, Rust, Go...) took decades to evolve comprehensive libraries of concurrent abstractions and data structures. In the last couple years, a handful of contributors to the OCaml ecosystem have been implementing basic libraries for concurrent and parallel programs in OCaml, in particular `Saturn` [Karvonen and Morel 2025], a library of lock-free thread-safe data structures (stacks, queues, a work-stealing queue, a skip list, a bag and a hash table), `Eio` [Madhavapeddy and Leonard 2025], a library of asynchronous IO and structured concurrency, and `Kcas` [Karvonen 2025a], a library offering a software-transactional-memory abstraction.

Concurrent algorithms and data structures are extremely difficult to reason about. Their implementations tend to be fairly short, a few dozens of lines. There is only a handful of experts able to write such code, and many potential users. They are difficult to test comprehensively. These characteristics make them ideally suited for mechanized program verification.

We embarked on a mission to mechanize correctness proofs of OCaml concurrent algorithms and data structures as they are being written, in contact with their authors, rather than years later. In the process, we not only gained confidence in these complex new building blocks, but we also improved the OCaml language and its verification ecosystem.

OCaml language features. When studying the new codebases of concurrent and parallel data structures, we found a variety of unsafe idioms, working around expressivity or performance limitations with the OCaml language support for lock-free concurrent data structures. In particular, the support for *atomic references* in the OCaml library proved inadequate, as idiomatic concurrent data-structures need the more expressive feature of *atomic record fields*. We designed an extension of OCaml with atomic record fields, implemented it as a an experimental compiler variant, and

succeeded in getting it integrated in the upstream OCaml compiler: it should be available as part of OCaml 5.4, which is not yet released at the time of writing.

Verification tools for concurrent programs. The state-of-the-art approach for mechanized verification of fine-grained concurrent algorithms is Iris [Jung, Krebbers, Jourdan, Bizjak, Birkedal and Dreyer 2018], a mechanized *higher-order* concurrent separation logic with *user-defined ghost state*. Its expressivity allows to precisely capture delicate invariants, and to reason about the linearization points of fine-grained concurrent algorithms, including external [Vindum, Frumin and Birkedal 2022] and future-dependent [Chang, Jung, Sharma, Tassarotti, Kaashoek and Zeldovich 2023; Jung, Lepigre, Parthasarathy, Rapoport, Timany, Dreyer and Jacobs 2020; Vindum and Birkedal 2021] linearization. Iris provides a generic mechanism to define programming languages and program logics for them. Most existing Iris concurrent verification work has been performed in HeapLang, the exemplar Iris language; a concurrent, imperative, untyped, call-by-value functional language.

To the best of our knowledge, it is currently the closest language to OCaml 5 in the Iris ecosystem — we review the existing frameworks in Section 12. We started our verification effort in HeapLang, but it eventually proved impractical to verify realistic OCaml libraries. Indeed, it lacks basic abstractions such as algebraic data types (tuples, mutable and immutable records, variants) and mutually recursive functions. Verifying OCaml programs in HeapLang requires difficult translation choices and introduces various encodings, to the point that the relation between the source and verified programs can become difficult to maintain and reason about. It also has very few standard data structures that can be directly reused. These limitations are well-known in the Iris community.

We created a new Iris language, ZooLang, that can better express concurrent OCaml programs. Its feature set grew over time as we applied it to more verification scenarios, and we now believe that it allows practical verification of fine-grained concurrent OCaml 5 programs — including the use of our atomic record fields which were co-designed with ZooLang. We were influenced by the Perennial framework [Chajed, Tassarotti, Kaashoek and Zeldovich 2019], which achieved similar goals for the Go language with a focus on crash-safety. As in Perennial, we also provide a translator from (a subset of) OCaml to ZooLang: ocaml2zoo. We start from OCaml code and call our translator to obtain a deep ZooLang embedding inside Rocq; we can use lightweight annotations to guide the translation. Inside Rocq we define specifications using Iris, and prove them correct with respect to the ZooLang version, which is syntactically very close to the original OCaml source. We call the resulting framework Zoo.

One notable current limitation of ZooLang is that it assumes a sequentially-consistent memory model, whereas OCaml offers a weaker memory model [Dolan, Sivaramakrishnan and Madhavapeddy 2018]. We wanted to ensure that we supported practical verification in a sequentially-consistent setting first; in the future we plan to adopt the OCaml memory model as formalized in Cosmo [Mével, Jourdan and Pottier 2020]. We discuss the impact of this difference in Section 3.5.

Specified OCaml semantics. Our Iris mechanization of ZooLang defines an operational semantics and a corresponding program logic. Our users on the other hand run their program through the standard OCaml implementation, which is not verified and does not have a precise formal specification. To bridge this formal-informal gap as well as reasonably possible, we carefully audit our ZooLang semantics to ensure that they coincide with OCaml's.

In doing so we discovered a hole in state-of-the-art language semantics for program verification (not just for OCaml), which is the treatment of *physical equality* (pointer equality). Physical equality is typically exposed to language users as an efficient but under-specified equality check, as the physical identity of objects may or may not be preserved by various compiler transformations. It is an essential aspect of concurrent programs, as it underlies the semantics of important atomic instructions such as compare_and_set. We found that the current informal semantics in OCaml is

incomplete, it does not allow to reason on programs that use structured data which mix mutable and immutable constructors. Existing formalizations of physical equality in verification frameworks typically restrict it to primitive datatypes, but idiomatic concurrent programs do not fit within this restriction. We propose a precise specification of physical equality in Zoo that scales to the verification of all the concurrent programs we encountered.

Worse, our discussions with the maintainers of the OCaml implementation showed that implementors guarantee weaker properties of physical equalities than users assume, in particular they may allow *unsharing*, which makes some existing concurrent programs incorrect. We propose a small new language feature for OCaml, per-constructor unsharing control, which we also integrate in our ZooLang translation, to fix affected programs and verify them. Finally, we discussed these subtleties with authors who axiomatize physical equality within Rocq for the purpose of efficient extraction, and we found out that some subtleties we discovered could translate into incorrectness in their axiomatization, requiring careful restrictions.

Verification results. We verified a small library for ZooLang, typically a subset of the OCaml standard library. It can serve as building blocks to define our concurrent data structures. (The lack of such a reusable standard library is a current limitation of HeapLang.) We verified a specific component of the Eio library, whose author Thomas Leonard had pointed to us as being delicate to reason about and worth mechanizing. Finally, we verified a large subset of the Saturn library. Several of these data structures contained verification challenges, which we will describe in the relevant section. The main Saturn concurrent structures remaining unverified are a skip-list and a hashtable; we have verified its work-stealing queue [Chase and Lev 2005], but do not discuss it here for space reasons.

Contributions. In summary, we claim the following contributions:

(1) Zoo, a practical program verification framework aimed at concurrent OCaml programs, mechanized in Rocq. The language ZooLang comes with a program logic expressed in the Iris concurrent separation logic. A translator ocaml2zoo generates Rocq embeddings from source OCaml programs, and works well with OCaml tooling (dune support).

(2) The verification (in a sequentially-consistent model) of important structures coming from Saturn, the OCaml 5 library of lock-free data structures. Our implementations and invariants sometimes improve over the Iris state of the art for those data structures.

(3) The extension of OCaml with atomic record fields, which after significant design, implementation and discussion work have now been integrated into upstream OCaml.

(4) The identification of blind spots in existing specifications of *physical equality*, and a new specification precise enough to reason about compare-and-set in concurrent programs. In the process we identified a potential bug in existing OCaml programs related to *unsharing*, and we propose a small language extension to let users selectively disable unsharing.

Artifact. We include an anonymous artifact containing (1) our experimental fork of the OCaml compiler implementing atomic arrays (Section 4.2) and generative constructors (Section 5.3), (2) the source of the ocaml2zoo translator, and (3) the Zoo repository, which includes ZooLang and its metatheory (theories/zoo), verified OCaml source (in lib/) and corresponding proofs (in theories/). In particular, the output of ocaml2zoo is included, for example mpmc_stack_1.ml in lib/zoo_saturn/ is translated into mpmc_stack_1__code.v in theories/zoo_saturn/.

2 Zoo in practice

The core of Zoo is ZooLang: a concurrent, imperative, untyped, functional programming language fully formalized in Rocq. Its semantics has been designed to match OCaml.

```
type 'a t = 'a list Atomic.t        Definition stack_create : val :=
let create () = Atomic.make []        fun: <> => ref [].

let rec push t v =                  Definition stack_push : val :=
  let old = Atomic.get t in           rec: "push" "t" "v" =>
  let new_ = v :: old in                let: "old" := !"t" in
  if not (Atomic.compare_and_set        let: "new_" := "v" :: "old" in
          t old new_)                   if: ~ CAS "t".[contents] "old" "new_"
  then (                                then (
    Domain.cpu_relax () ;                 domain_yield () ;;
    push t v                              "push" "t" "v"
  )                                     ).

let rec pop t =                     Definition stack_pop : val :=
  match Atomic.get t with             rec: "pop" "t" =>
  | [] -> None                          match: !"t" with
  | v :: new_ as old ->                 | [] => §None
      if Atomic.compare_and_set         | ("v" :: "new_") as "old" =>
          t old new_                        if: CAS "t".[contents] "old" "new_"
      then Some v                         then 'Some( "v" )
      else (                              else (
        Domain.cpu_relax () ;               domain_yield () ;;
        pop t                               "pop" "t"
      )                                   )
                                        end.
```

Fig. 1. A concurrent stack in OCaml and its automatic ZooLang translation

ZooLang comes with a program logic based on Iris, proved correct with respect to its small-step operational semantics. The reasoning rules are expressed in separation logic (including rules for the different constructs of the language) along with Rocq tactics that integrate into the Iris proof mode [Krebbers, Jourdan, Jung, Tassarotti, Kaiser, Timany, Charguéraud and Dreyer 2018]. In addition, it supports Diaframe [Mulder, Krebbers and Geuvers 2022], enabling proof automation.

2.1 Translation from OCaml to ZooLang

While ZooLang lives in Rocq, we want to verify OCaml programs. To connect them we provide the tool ocaml2zoo to translate OCaml source files[1] into Rocq files containing ZooLang code. This tool can process entire dune projects, and supports library dependencies.

The supported OCaml fragment includes: tuples, variants, records and inline records, shallow **match**, atomic record fields, unboxed types, toplevel mutually recursive functions.

In Figure 1 we include the OCaml implementation of a simple lock-free concurrent stack [Treiber 1986], and its automatic translation to ZooLang, demonstrating that readability is preserved. Readability is important as users constantly see program fragments during interactive verification.

Rocq term	t		
constructor	C		
projection	$proj$		
record field	fld		
identifier	s, f	$\in$	String
integer	n	$\in$	$\mathbb{Z}$
boolean	b	$\in$	$\mathbb{B}$
binder	x	::=	`<>` $\mid s$
unary operator	$\oplus$	::=	`~` $\mid$ `-`
binary operator	$\otimes$	::=	`+` $\mid$ `-` $\mid$ `*` $\mid$ `‘quot‘` $\mid$ `‘rem‘` $\mid$ `‘land‘` $\mid$ `‘lor‘` $\mid$ `‘lsl‘` $\mid$ `‘lsr‘`
			$\mid$ `<=` $\mid$ `<` $\mid$ `>=` $\mid$ `>` $\mid$ `=` $\mid$ `≠` $\mid$ `==` $\mid$ `!=`
			$\mid$ `and` $\mid$ `or`
expression	e	::=	$t \mid s \mid$ `#`$n \mid$ `#`b
			$\mid$ `fun:` $x_1 \ldots x_n$ `=>` $e \mid$ `rec:` $f\ x_1 \ldots x_n$ `=>` $e \mid e_1\ e_2$
			$\mid$ `let:` x `:=` e_1 `in` $e_2 \mid e_1$ `;;` e_2
			$\mid$ `let:` $f\ x_1 \ldots x_n$ `:=` e_1 `in` $e_2 \mid$ `letrec:` $f\ x_1 \ldots x_n$ `:=` e_1 `in` e_2
			$\mid$ `let:` `‘`$C\ x_1 \ldots x_n$ `:=` e_1 `in` $e_2 \mid$ `let:` $x_1, \ldots, x_n$ `:=` e_1 `in` e_2
			$\mid \oplus e \mid e_1 \otimes e_2$
			$\mid$ `if:` e_0 `then` e_1 `(else` $e_2)^?$
			$\mid$ `for:` x `:=` e_1 `to` e_2 `begin` e_3 `end`
			$\mid$ `§`$C \mid$ `‘`$C\ (e_1, \ldots, e_n) \mid (e_1, \ldots, e_n) \mid e$`.<`$proj$`>`
			$\mid$ `[]` $\mid e_1$ `::` e_2
			$\mid$ `‘`$C\ \{e_1, \ldots, e_n\} \mid \{e_1, \ldots, e_n\} \mid e$`.{`$fld$`}` $\mid e_1$ `<-{`fld`}` e_2
			$\mid$ `ref` $e \mid$ `!`$e \mid e_1$ `<-` e_2
			$\mid$ `match:` e_0 `with` $br_1 \mid \ldots \mid br_n$ `(`$\mid$ `_ (as` $s)^?$ `=>` $e)^?$ `end`
			$\mid e$`.[`fld`]` $\mid$ `Xchg` $e_1\ e_2 \mid$ `CAS` $e_1\ e_2\ e_3 \mid$ `FAA` $e_1\ e_2$
			$\mid$ `Proph` $\mid$ `Resolve` $e_0\ e_1\ e_2$
branch	br	::=	$C\ (x_1 \ldots x_n)^?\ ($`as` $s)^?$ `=>` e
			$\mid$ `[]` `(as` $s)^?$ `=>` $e \mid x_1$ `::` x_2 `(as` $s)^?$ `=>` e
toplevel value	v	::=	$t \mid$ `#`$n \mid$ `#`b
			$\mid$ `fun:` $x_1 \ldots x_n$ `=>` $e \mid$ `rec:` $f\ x_1 \ldots x_n$ `=>` e
			$\mid$ `§`$C \mid$ `‘`$C\ (v_1, \ldots, v_n) \mid (v_1, \ldots, v_n)$
			$\mid$ `[]` $\mid v_1$ `::` v_2

Fig. 2. ZooLang syntax (omitting mutually recursive toplevel functions)

2.2 The full language

The ZooLang syntax is given in Figure 2[2], omitting mutually recursive toplevel functions that are treated specially. Expressions include standard constructs like booleans, integers, anonymous functions (that may be recursive), applications, **let** bindings, sequence, unary and binary operators, conditionals, **for** loops, tuples. In any expression, one can refer to a Rocq term representing a ZooLang value (of type `val`) using its Rocq identifier. ZooLang is deeply embedded: variables (bound by functions and **let**) are quoted as strings.

Data constructors (immutable memory blocks) are supported: `§`C represents a constant constructor (*e.g.* `§None`), and `‘`$C\ (e_1, \ldots, e_n)$ represents a non-constant constructor (*e.g.* `‘Some(` e `)`).

---

[1]Actually, ocaml2zoo processes binary annotation files (`.cmt` files).

[2]More precisely, it is the syntax of the surface language, including Rocq notations.

Projections of the form $e.\langle proj \rangle$ (including on tuples: $(x, y).\langle 1 \rangle$) can be used to obtain a specific component of a tuple or data constructor. ZooLang supports shallow pattern matching (patterns cannot be nested) on data constructors with an optional fallback case.

Mutable memory blocks are constructed using either the untagged record syntax $\{e_1, \ldots, e_n\}$ or the tagged record syntax `C $\{e_1, \ldots, e_n\}$. Reading a record field can be performed using $e.\{fld\}$ and writing to a record field using $e_1 \leftarrow \{fld\}\ e_2$. Pattern matching can also be used on mutable tagged blocks provided that cases do not bind anything — in other words, only the tag is examined, no memory access is performed. Mutable references are supported: ref e creates a reference, $!e$ reads a reference and $e_1 \leftarrow e_2$ writes into a reference. There is no built-in syntax for arrays; they are supported through the **Array** standard library module (*e.g.* array_make).

Note that ZooLang follows OCaml in sometimes eschewing orthogonality to provide more compact memory representations: constructors are *n*-ary instead of taking a tuple as parameter, and the tagged record syntax is distinct from a constructor taking a mutable record as parameter. In each case the simplifying encoding would introduce an extra indirection in memory, which is absent from the ZooLang semantics. Performance-conscious experts care about these representation choices, and we care about faithfully modeling their programs.

Parallelism is mainly supported through the **Domain** standard library module, including domain-local storage. Atomic operations are provided as built-in constructs (Xchg, CAS, FAA; see Section 3.3).

The Proph and Resolve constructs model *prophecy variables* [Jung, Lepigre, Parthasarathy, Rapoport, Timany, Dreyer and Jacobs 2020], see Section 3.4.

2.3 Specifications and proofs

Once the translation to ZooLang is done, the user can write specifications and prove them in Iris. For instance, the specification of the stack_push function from Figure 1 could be:

```
Lemma stack_push_spec t ι v :
  <<< stack_inv t ι
    | ∀∀ vs, stack_model t vs >>>
    stack_push t v @ ↑ι
  <<< stack_model t (v :: vs)
    | RET (); True >>>.
```

It uses a *logically atomic specification* [da Rocha Pinto, Dinsdale-Young and Gardner 2014], which has been proven [Birkedal, Dinsdale-Young, Guéneau, Jaber, Svendsen and Tzevelekos 2021] to be equivalent to *linearizability* [Herlihy and Wing 1990] in sequentially consistent memory models.

Similarly to Hoare triples, the specification is formed of a precondition and a postcondition, represented in angle brackets. But each is split in two parts, a *public* or *atomic* condition, and a *private* condition. Following standard Iris notations, the private conditions are on the outside (first line of the precondition, last line of the postcondition) and the atomic conditions are inside.

For this particular operation, the private postcondition is trivial. The private precondition stack_inv t is the stack invariant. Intuitively, it asserts that t is a valid concurrent stack. More precisely, it defines a concurrent protocol that t must respect at all times.

The atomic pre- and post-conditions specify the linearization point of the operation: during the execution of stack_push, the abstract state of the stack held by stack_model is atomically updated from *vs* to *v :: vs* when *v* is atomically pushed at the top of the stack.

3 Zoo features

In this section, we review the salient features of Zoo, which we found lacking when we attempted to use HeapLang to verify real-world OCaml programs. Providing a better Iris language than

HeapLang for this problem domain is a contribution of our work; others are welcome to use ZooLang for their own verification effort, or to reuse our designs by integrating specific features into their own Iris language.

3.1 Algebraic data types

ZooLang is an untyped language but, to write interesting programs, it is convenient to work with abstractions like algebraic data types. To simulate tuples, variants and records, we designed a machinery to define projections, constructors and record fields.

For example, one may define a list-like type with:

```
Notation "'Nil'"  := (in_type "t" 0) (in custom zoo_tag).
Notation "'Cons'" := (in_type "t" 1) (in custom zoo_tag).
```

Users do not need to write this incantation directly, as they are generated by ocaml2zoo from the OCaml type declarations. Suffice it to say that it introduces the two tags in the zoo_tag custom entry, on which the notations for data constructors rely. The in_type term is needed to distinguish the tags of distinct data types; crucially, it cannot be simplified away by Rocq, as this could lead to confusion during the reduction of expressions.

One may then directly use the tags Nil and Cons in data constructors:

```
Definition map : val := rec: "map" "fn" "t" =>
  match: "t" with
  | Nil => §Nil
  | Cons "x" "t" =>
      let: "y" := "fn" "x" in
      ‘Cons( "y", "map" "fn" "t" )
  end.
```

Similarly, one can define a record-like type with two mutable fields f1 and f2:

```
Notation "'f1'" := (in_type "t" 0) (in custom zoo_field).
Notation "'f2'" := (in_type "t" 1) (in custom zoo_field).

Definition swap : val := fun: "t" =>
  let: "f1" := "t".{f1} in
  "t" <-{f1} "t".{f2} ;; "t" <-{f2} "f1".
```

3.2 Mutually recursive functions

Zoo supports non-recursive (fun: $x_1 \ldots x_n$ => e) and recursive (rec: $f\ x_1 \ldots x_n$ => e) functions but only *toplevel* mutually recursive functions. It is non-trivial to properly handle mutual recursion: when applying a mutually recursive function, a naive approach would replace calls to sibling functions by their respective bodies, but this typically makes the resulting expression unreadable and can create proof-checking performance issues during verification. To prevent it, the mutually recursive functions have to know one another to preserve their names during β-reduction. For instance, one may define two mutually recursive functions f and g as follows. ocaml2zoo generates some additional boilerplate to control the recursive unfolding.

```
Definition f_g := (
  recs: "f" "x" => "g" "x"
  and:  "g" "x" => "f" "x"
)%zoo_recs.
```

3.3 Fine-grained concurrent primitives

Zoo supports concurrent primitives both on atomic references (from **Atomic**) and atomic record fields (from **Atomic.Loc**, see Section 4.1) according to the table below. The OCaml expressions listed on the left translate into the Zoo expressions on the right.

OCaml	Zoo
Atomic.get e	!e
Atomic.set e_1 e_2	e_1 <- e_2
Atomic.exchange e_1 e_2	Xchg e_1.[contents] e_2
Atomic.compare_and_set e_1 e_2 e_3	CAS e_1.[contents] e_2 e_3
Atomic.fetch_and_add e_1 e_2	FAA e_1.[contents] e_2
Atomic.Loc.exchange [%atomic.loc e_1.f] e_2	Xchg e_1.[f] e_2
Atomic.Loc.compare_and_set [%atomic.loc e_1.f] e_2 e_3	CAS e_1.[f] e_2 e_3
Atomic.Loc.fetch_and_add [%atomic.loc e_1.f] e_2	FAA e_1.[f] e_2

One notable aspect of this translation is that atomic accesses (**Atomic**.get and **Atomic**.set) correspond to plain loads and stores. This is because we are working in a sequentially consistent memory model: there is no difference between atomic and non-atomic memory locations.

3.4 Prophecy variables

Lock-free algorithms exhibit complex behaviors. To tackle them, Iris provides powerful mechanisms such as *prophecy variables* [Jung, Lepigre, Parthasarathy, Rapoport, Timany, Dreyer and Jacobs 2020]. Essentially, prophecy variables can be used to predict the future of the program execution and reason about it. They are key to handle *future-dependent linearization points*: linearization points that may or may not occur at a given location in the code depending on a future observation.

Zoo supports prophecy variables through the Proph and Resolve expressions — as in Heaplang, the canonical Iris language. In OCaml, these expressions correspond to **Zoo**.proph and **Zoo**.resolve, that are recognized by ocaml2zoo.

3.5 Future work: the OCaml memory model

Our current formalization of Zoo assumes sequential consistency; this does not faithfully model all possible behaviors of concurrent OCaml programs, which uses a more relaxed memory model [Dolan, Sivaramakrishnan and Madhavapeddy 2018]. Some concurrent algorithms, such as the Treiber stack, do not contain any data races between atomic and non-atomic locations, so their behavior is identical between both memory models. (Our specifications do not cover the synchronization guarantees on user-provided data, which varies between implementations in relaxed models.). Other algorithms do contain data races, and our formal correctness result must be taken with the caveat that it does not describe all observable behaviors in the actual OCaml program.

We made the choice to focus on proving correctness of these subtle concurrent algorithms in the simpler setting of sequential consistency first, to encounter and solve the obstacles to practical verification of concurrent programs. We intend to migrate Zoo to the OCaml memory model as formalized by Mével, Jourdan and Pottier [2020], and have started preliminary work in this direction, introducing its *views* in a work-in-progress version of the ZooLang program logic.

4 OCaml extensions for fine-grained concurrent programming

Over the course of this work, we studied efficient fine-grained concurrent OCaml programs written by experts. This revealed various limitations of OCaml in these domains, that those experts would work around using unsafe casts, often at the cost of both readability and memory-safety; and also some mismatches between their mental model of the semantics of OCaml and the mental model

393 used by the OCaml compiler authors. We worked on improving OCaml itself to reduce these
394 work-arounds or semantic mismatches.

4.1 Atomic record fields

OCaml 5 offers a type `'a Atomic.t` of atomic references exposing sequentially-consistent atomic
operations. Data races on non-atomic mutable locations has a much weaker semantics and is gen-
erally considered a programming error. For example, the Michael-Scott concurrent queue [Michael
and Scott 1996] relies on a linked list structure that could be defined as follows:

```
type 'a node = Nil | Cons of { value : 'a; next : 'a node Atomic.t }
```

Performance-minded concurrency experts dislike this representation, because `'a Atomic.t`
introduces an indirection in memory: it is represented as a pointer to a block containing the value
of type `'a`. Instead, they use something like the following:

```
type 'a node = Nil | Cons of { mutable next: 'a node; value: 'a }
let as_atomic : 'a node -> 'a node Atomic.t option = function
  | Nil -> None
  | (Cons _) as record -> Some (Obj.magic record : 'a node Atomic.t)
```

Notice that the `next` field of the `Cons` constructor has been moved first in the type decla-
ration. Because the OCaml compiler respects field-declaration order in data layout, a value
`Cons { next; value }` has a similar low-level representation to a reference (atomic or not) point-
ing at `next`, with an extra argument. The code uses `Obj.magic` to unsafely cast this value to an
atomic reference, which appears to work as intended.

`Obj.magic` is a shunned unsafe cast (the OCaml equivalent of `unsafe` or `unsafePerformIO`).
It is very difficult to be confident about its usage given that it may typically violate assumptions
made by the OCaml compiler and optimizer. In the example above, casting a two-fields record into
a one-argument atomic reference may or may not be sound — but it gives measurable performance
improvements on concurrent queue benchmarks.

It is possible to statically forbid passing `Nil` to `as_atomic` to avoid error handling, by turning
`'a node` into a GADT indexed over a type-level representation of its head constructor. This pattern
can be found in the Kcas [Karvonen 2025a] library by Vesa Karvonen. It is difficult to write and use
correctly, in particular as unsafe casts can sometimes hide errors in the intended static discipline.

Note that this unsafe approach only works for the first field of a record, so it is not applicable to
records that hold several atomic fields, such as the toplevel record storing atomic `front` and `back`
pointers for the concurrent queue.

4.1.1 Our atomic fields proposal. In May-June 2024 we proposed a design for atomic record fields as
an OCaml language change proposal . Declaring a record field atomic simply requires an `[@atomic]`
attribute — and could eventually become a proper keyword of the language.

```
(* re-implementation of atomic references *)
type 'a atomic_ref = { mutable contents : 'a [@atomic]; }

(* concurrent linked list *)
type 'a node = Nil | Cons of { value: 'a; mutable next : 'a node [@atomic]; }

(* bounded SPSC circular buffer *)
type 'a bag = { data : 'a Atomic.t array;
                mutable front: int [@atomic];
                mutable back: int [@atomic]; }
```

The design difficulty is to express atomic operations on atomic record fields. For example, if buf has type 'a bag above, then one naturally expects the existing notation buf.front to perform an atomic read and buf.front <- n to perform an atomic write. But how would one express exchange, compare-and-set and fetch-and-add? We would like to avoid adding a new primitive language construct for each atomic operation.

Two different designs have been proposed for atomic operations on atomic record fields.

(1) The original, "full" design in our proposal introduces a built-in type ('r, 'a) **Atomic.Field**.t that denotes a field/index of type 'a within a record of type 'r, with a syntax extension [%atomic.loc <field>] to construct such field description, and atomic primitives in a module **Atomic.Field**, that need both the record value of type 'r and the field description. For example, the standard library exposes:

```
val Atomic.Field.fetch_and_add : 'a -> ('a, int) Atomic.Field.t -> int -> int
```
and users can write:
```
let preincrement_front (buf : 'a bag) : int =
  Atomic.Loc.fetch_and_add buf [%atomic.field front] 1
```
where [%atomic.field front] has type ('a bag, int) **Atomic.Loc**.t. Internally, a value of type ('r, 'a) **Atomic.Field**.t is just an integer offset locating the field within the record.

(2) An alternative "simple" design, which was proposed by Basile Clément, introduces a built-in type 'a **Atomic.Loc**.t for an atomic location that holds an element of type 'a, with a syntax extension [%atomic.loc <expr>.<field>] to construct such locations. Atomic primitives operate on values of type 'a **Atomic.Loc**.t and they are exposed as functions of the module **Atomic.Loc**. For example, the standard library exposes:

```
val Atomic.Loc.fetch_and_add : int Atomic.Loc.t -> int -> int
```
and users can write:
```
let preincrement_front (buf : 'a bag) : int =
  Atomic.Loc.fetch_and_add [%atomic.loc buf.front] 1
```
where [%atomic.loc buf.front] has type int **Atomic.Loc**.t. Internally, a value of type 'a **Atomic.Loc**.t can be represented as a pair of a record and an integer offset for the desired field, and the atomic.loc construction builds this pair in a well-typed manner. When a primitive of the **Atomic.Loc** module is applied to an atomic.loc expression, the compiler can optimize away the construction of the pair — but it would happen if there was an abstraction barrier between the construction and its use.

The simple design has the strong advantage of exposing simpler types to users. The full design involves more complex types, but a simpler data representation that does not rely on specific compiler optimizations to generate efficient code, even across abstraction boundaries. As remarked by Leo White, the simpler type 'a **Atomic.Loc**.t can be reconstructed as a dependent pair of a 'r and a ('r, 'a) **Atomic.Field**.t, which is expressible in OCAML as a GADT:

```
type 'a loc = Loc : 'r * ('r, 'a) Atomic.Field.t -> 'a loc
```

In August and September 2024 we implemented the simple design as an experimental version of the OCAML compiler. We got a full code review by Olivier Nicole. When we approached the OCAML maintainers to integrate the feature upstream, they asked us to reconsider implementing the full design, based on feedback in this direction by Vesa Karvonen (author of Kcas and Saturn).

We implemented the full design between November 2024 and January 2025. We found that it is harder to implement in the type-checker. The extension form [%atomic.loc buf.front] of the simple design has typing rules that are very similar to a field access buf.front. On the other hand,

[%atomic.loc front] interacts in a non-trivial way with the OCaml machinery for type-based disambiguation of record fields — when several records exist with a field named front.

For technical reasons, there is also a non-trivial interaction with the type-checking of inline record types (record types that are not defined by themselves but only as the argument of a sum type constructor), which currently prevents from using this approach with inline records, who are unfortunately common in efficient mutable data structures. In February and March 2025 we worked with OCaml maintainers experts to try to lift these limitations and proposed implementation changes for inline records, but we failed to reach a consensus on the best design.

In April and May, we managed to build consensus among OCaml maintainers that the simple design, with a simpler user-facing design and a simpler implementation, was a good compromise in view of the significant difficulties of the full design with inline records. Our implementation of the simple design was reviewed again, we wrote user documentation, and it was integrated upstream in May 2025, to be included in the upcoming release of OCaml 5.4.

Limitation: no support for cache contention. The type **Atomic**.t comes with a function

```
val Atomic.make_contended : 'a -> 'a Atomic.t
```

that ensures that the returned atomic reference is allocated with enough alignment and padding to sit alone on its cache line, to avoid performance issues caused by false sharing. Currently there is no such support for padding of atomic record fields (we are planning to work on this), so the less-compact atomic references remain preferable in certain scenarios.

4.2 Atomic arrays

On top of our atomic record fields, we have implemented experimental support for atomic arrays, another facility commonly requested by authors of efficient concurrent programs. Our previous example of a concurrent bag of type 'a bag used a backing array of type 'a **Atomic**.t **array**, which contains more indirections than may be desirable, as each array element is a pointer to a block containing the value of type 'a, instead of storing the value of type 'a directly in the array.

Our implementation of atomic arrays builds on top of the type 'a **Atomic.Loc**.t we described in the previous section, and it relies on two new low-level primitives provided by the compiler:

```
val Atomic_array.index : 'a array -> int -> 'a Atomic.Loc.t
val Atomic_array.unsafe_index : 'a array -> int -> 'a Atomic.Loc.t
```

The function index takes an array and an integer index within the array, and returns an atomic location into the corresponding element after performing a bound check. unsafe_index omits the boundcheck — additional performance at the cost of memory-safety — and allows to express the atomic counterpart of the unsafe operations **Array**.unsafe_get and **Array**.unsafe_set. The atomic primitives of the module **Atomic.Loc** can then be used on these indices; our implementation implements a library module on top of these primitives to provide a higher-level layer to the user, with direct array operations such as:

```
val Atomic_array.exchange : 'a Atomic_array.t -> int -> 'a -> 'a
val Atomic_array.unsafe_exchange : 'a Atomic_array.t -> int -> 'a -> 'a
```

5 Physical equality

The notion of *physical equality* is ubiquitous in fine-grained concurrent algorithms. It appears not only in the semantics of the (==) operator, but also in the semantics of the **Atomic**.compare_and_set primitive, which atomically sets an atomic reference to a desired value if its current content is physically equal to an expected value. This primitive is commonly used to try committing an atomic operation in a retry loop, as in the push and pop functions of Figure 1.

5.1 Physical equality in HEAPLANG

In HEAPLANG, this primitive is provided but restricted. Indeed, its semantics is only defined if either the expected or the desired value fits in a single memory word in the HEAPLANG value representation: literals (booleans, integers and pointers[3]) and literal injections[4]; otherwise, the program is stuck. In practice, this restriction forces the programmer to introduce an indirection [Iris development team 2025; Jung, Lepigre, Parthasarathy, Rapoport, Timany, Dreyer and Jacobs 2020; Vindum and Birkedal 2021] to physically compare complex values, *e.g.* lists. Furthermore, when the semantics is defined, values are compared using their ROCQ representations; physical equality boils down to ROCQ equality.

5.2 Physical equality in OCAML

In OCAML, *physical* equality is tricky. *Structural* equality v1 = v2, which we describe in Section 6, would be the preferred way of comparing values, and using physical equality v1 == v2 is often an unintentional mistake. However, physical equality is typically much faster than structural equality, as it compiles to only one assembly instruction instead of traversing the value. Also, the `Atomic.compare_and_set` requires the comparison to be atomic, ruling out structural equality.

Physical equality is in a counter-intuitive situation: it is very simple to *implement* (in the OCAML compiler, or in an interpreter, etc.) but difficult to *specify* precisely. To make verification practical, we need a specification at the level of *source* OCAML (or ZOO) programs, using a high-level representation of values, as close to the source as reasonably possible, that we call *abstract* values. On the other hand, its implementation typically work with *low-level* values, and its observable behavior depends on compiler transformations that happen in-between the two abstraction levels. This difficulty can result in dangerous gaps between the programming language used to write code and the semantics used for its verification.

ZOOLANG has a grammar of values, and most operations are specified by defining how they compute with ZOOLANG values. Its definition may look as follows in ROCQ (simplified):

```
Inductive literal :=                Inductive val :=
| Bool (b : bool)                   | Lit (lit : literal)
| Int (n : nat)                     | Recs (i : nat)
| Loc (l : location)                       (recs : list (binder * binder * expr))
| Proph (pid : prophet_id)          | Block (tag : nat) (vs : list val).
| Poison.
```

The value 'Cons(42, §Nil) is represented in ROCQ as Block 1 [Lit (Int 42), Block 0 []]. Notice that immutable blocks are represented in ROCQ using the Block constructor directly, and *not* as a location (Loc) allocated on the heap. We use locations only for *mutable* records. We would say that our representation of ZOOLANG values is high-level or *abstract*, as close to the surface syntax as reasonably possible. This distinction is important to make verification pleasant in practice, by reducing the number of locations and heap indirections that the programmer needs to work with during verification. A ZOOLANG tuple is directly a tuple, etc., and this design decision of using high-level values is important to the verification experience — in addition, assuming full ownership of arguments of immutable blocks would be incorrect.

It is tempting to specify, as HEAPLANG does, that physical equality decides equality between abstract values. This specification makes sense for immediate values (integers, booleans, constant constructors), and for mutable records which are compared by location. But it is incorrect on

---

[3]HEAPLANG allows arbitrary pointer arithmetic and therefore inner pointers. This is forbidden in both OCAML and ZOOLANG, as any reachable value has to be compatible with the garbage collector.

[4]HEAPLANG has no primitive notion of constructor, only pairs and injections (left and right).

immutable blocks, and HEAPLANG essentially does not specify its behavior on those values. Yet programmers use physical equality on immutable blocks in practice, as in our example of a Treiber stack of Figure 1.

Defining physical equality as ROCQ equality of abstract values is problematic in opposite ways:

(1) Some distinct abstract values are physically equal in OCaml, for example 0 and false. Their type differ, but it possible to store them in an existential type where they can be compared for physical equality:

```
type any = Any : 'a -> any
let test1 = Any false == Any 0 (* may return true *)
```

This shows that even on immediate values, specifying physical equality as equality of abstract values is convenient but incorrect in practice.

(2) A deeper problem is that some definitionally equal abstract values may be physically distinct. Consider for example the case of *immutable blocks* representing constructors and immutable records (as opposed to *mutable blocks* representing mutable records), *e.g.* Some 0. The physical comparison of Some 0 and Some 0 may return either true or false: we cannot determine the result of physical comparison just by looking at the abstract values.

To solve these problems we treat physical equality on abstract values is *non-deterministic* — even though the comparison instruction that implements it on low-level values is perfectly deterministic. The question is then: what guarantees do we get when physical equality returns true and when it returns false? Given such guarantees, denoted by val_physeq and val_physneq, the non-deterministic semantics is reflected in the logic through the following specification:

```
Lemma physeq_spec v1 v2 :
  {{{ True }}}
    v1 == v2
  {{{ b, RET #b; ⌜(if b then val_physeq else val_physneq) v1 v2⌝ }}}
```

The OCaml manual documents a partial specification for physical equality, which is precise for basic types such as integers or integer references, but does not clearly extend to structured values containing a mix of immutable and mutable constructors, which are present in the programs we verify. The only guarantee that it provides for all values is: if two values are physically equal, they are also structurally equal. For values that contain immutable constructors, we do not learn anything when they are physically distinct.

We will now explore the specifications of the true and false return cases. We describe our program verification requirements to suggest a precise enough semantics: if it does not say enough, we cannot prove our programs correct. In the other direction we describe some optimizations that OCaml implementations may perform (gathered through our discussion with OCaml maintainers), in particular the current compiler, that may rule out certain stronger specifications as incorrect.

Remark. It is tempting to state that physical equality implies equality of ROCQ representations, but incorrect in general as we have seen. Existing work on the modelling of OCaml physical equality within proof assistants have typically made this simplification, restricting the set of supported values to preserve soundness. We discussed our examples with authors of such earlier formalizations, and this has sometimes uncovered soundness issues on their side, as we discuss in Section 12.3.

5.3 When physical equality returns true

Let us go back to the concurrent stack of Figure 1 and more specifically the push function. Its atomic specification, given in Section 2.3, states that if we push a value v onto a stack whose current model is vs, then it atomically becomes a stack of model $v :: vs$. To prove this specification we rely

on the fact that, if `Atomic`.`compare_and_set` returns `true`, the current list must be equal to *vs*, in the sense of Rocq equality for Zoo values. This equality is strictly stronger than structural equality on mutable types, so we need a more precise specification than provided by the OCaml manual.

Zoo supports the following fragment of OCaml values: booleans, integers, mutable blocks (pointers), immutable blocks, functions.

The easy cases are mutable blocks and functions. Each of these two classes is disjoint from all others. We can reasonably assume that, when physical equality returns `true` and one of the compared values belongs to either of these classes, the two values are actually the same in Rocq. As far as we are aware, there is no optimization that could break this.

Booleans, integers and empty immutable blocks (constant constructors) are all represented as immediate integers in OCaml's low-level representation. This encoding induces conflicts: two abstract values that are distinct in Rocq may have the same low-level representation. The semantics of unrestricted physical equality has to reflect these conflicts: on those values that have an immediate representation, our specification does not state that physical equality includes equality of abstract values, it introduces a (simplified) notion of low-level representation and only states that those representations are equal.

Finally, let us consider the case of non-empty immutable blocks. At runtime, they are represented by pointers to tagged memory blocks. At first approximation, it is tempting to say that physically equal immutable blocks really are definitionally equal in Rocq. Alas, this is not true. To explain why, we have to recall that the OCaml compiler may perform *sharing*: immutable blocks containing physically equal fields may be shared. For example, the following tests may return `true`:

```
let test1 = Some 0 == Some 0 (* true *)
let test2 = [0;1]  == [0;1]  (* true *)
```

On its own, sharing is not a problem. However, coupled with representation conflicts, it can be surprising. Indeed, consider the any type we introduced previously:

```
type any = Any : 'a -> any
```

The following tests may return `true` (they do with `ocamlopt`, not `ocamlc`):

```
let test1 = Any false == Any 0 (* true *)
let test2 = Any None  == Any 0 (* true *)
let test3 = Any []    == Any 0 (* true *)
```

Now, going back to the push function of Figure 1, we have a problem. Given a stack of any, it is possible for the `Atomic`.`compare_and_set` to observe a current list (*e.g.* the one-element list [`Any` 0]) physically equal to the expected list (*e.g.*, [`Any` `false`]) while these are actually distinct in Rocq. In short, the expected specification of Section 2.3 is incorrect: we may not get *v* :: *vs* back in the model, but a list *v* :: *vs'* where *vs'* is physically equal to *vs* but not the same abstract value. To fix this discrepancy, we would need to weaken all our specifications to be formulated *modulo physical equality*, which is non-standard and quite burdensome.

We believe this really is a shortcoming, at least from the verification perspective. Therefore, we extended OCaml with *generative immutable blocks* . These generative blocks are just like regular immutable blocks, except they cannot be shared. Hence, if physical equality on two generative blocks returns `true`, these blocks are definitionally equal in Rocq. At user level, this notion is materialized by *generative constructors*. For instance, to verify the expected push specification, we can use a generative version of lists:

```
type 'a glist =
  | Nil
  | Cons of 'a * 'a glist [@generative]
```

```
687   type state =                                    type t =
688     | Open of Unix.file_descr                       { mutable ops: int [@atomic];
689     | Closing of (unit -> unit)                        mutable state: state [@atomic]; }
690
691   let make fd = { ops = 0; state = Open fd }
692
693   let closed = Closing (fun () -> ())
694   let close t =
695     match t.state with
696     | Closing _ -> false
697     | Open fd as prev ->
698         let close () = Unix.close fd in
699         let next = Closing close in
700         if Atomic.Loc.compare_and_set [%atomic.loc t.state] prev next then (
701           if t.ops == 0
702           && Atomic.Loc.compare_and_set [%atomic.loc t.state] next closed
703           then close () ;
704           true
705         ) else false
```

Fig. 3. **Rcfd** module from Eio (excerpt)

We modified the Zoo translator to support those generative constructors, and modified our implementation of Treiber stack to use the type 'a glist instead of 'a **list**, so that we could finally prove the expected, convenient specification. The [@generative] attribute is ignored by the standard OCaml compiler, and we have an experimental version that disables sharing for generative constructors.

5.4   When physical equality returns **false**

Most formalizations of physical equality in the literature do not give any guarantee when physical equality returns **false**. Many use-cases of physical equality, in particular retry loops, can be verified with only sufficient conditions on **true**. However, in some specific cases, more information is needed.

Consider the **Rcfd** module from the Eio library, an excerpt of which is given in Figure 3[5]. Thomas Leonard, its author, suggested that we verify this real-life example because of its intricate logical state. However, we found out that it is also relevant regarding the semantics of physical equality. Essentially, it consists in wrapping a file descriptor in a thread-safe way using reference-counting. At creation in the make function, the wrapper starts in the **Open** state. At some point, it can switch to the **Closing** state in the close function, and will remain **Closing** forever. Crucially, the **Open** state changes at most once to **Closing**, never to another **Open**.

The interest of **Rcfd** lies in the close function. First, the function reads the state. If this state is **Closing**, it returns **false**; the wrapper has been closed. If this state is **Open**, it tries to switch to the **Closing** state using **Atomic.Loc**.compare_and_set; if this attempt fails, it also returns **false**. In this particular case, we would like to prove that the wrapper has been closed, or equivalently that **Atomic.Loc**.compare_and_set cannot have observed **Open**. Intuitively, if we observed a different value then it must be **Closing**.

---

[5]We make use of *atomic record fields* as introduced in Section 4.1.

Obviously, we need some kind of guarantee related to the *physical identity* of **Open** when `Atomic.Loc`.compare_and_set returns `false`. If **Open** were a mutable block, we could argue that this block cannot be physically distinct from itself; no optimization we know of would allow that. Unfortunately, it is an immutable block, and immutable blocks are subject to more optimizations. In fact, something surprising but allowed[6] by OCaml can happen: *unsharing*, the dual of sharing. Indeed, any immutable block can be unshared, that is reallocated as if its definition was inlined. For example, the following test may theoretically return `false`:

```
let x = Some 0
let test = x == x (* false *)
```

Going back to **Rcfd**, we have a problem: in the second branch, the **Open** block corresponding to prev could be unshared, which would make `Atomic.Loc`.compare_and_set fail. Hence, we cannot prove the expected specification; in fact, the program as it is written has a bug.

To remedy this unfortunate situation, we propose to reuse the notions of generative immutable blocks, that we introduced to prevent sharing, to also forbid unsharing by the OCaml compiler — we implemented this in our experiment compiler branch.

Supporting this requires enriching the Zoo semantics so that each generative block is annotated with a *logical identifier*[7] representing its physical identity, much like a pointer for a mutable block. If physical equality on two generative blocks returns `false`, the two identifiers are necessarily distinct. Given this semantics, we can verify the close function. Indeed, if `Atomic.Loc`.compare_and_set fails, we now know that the identifiers of the two blocks, if any, are distinct. As the concurrent protocol has only one **Open** block whose identifier does not change, it cannot be the case that the current state is **Open**, hence it is **Closing**. We can verify this function after adding the following annotation:

```
type state =
  | Open of Unix.file_descr [@generative]
  | Closing of (unit -> unit)
```

5.5 Summary

In summary, we extended our abstract values with generative immutable blocks, and give the following specification to physical equality in ZooLang. It can also serves as a precise specification of physical equality of a practical fragment of OCaml:

- On values whose low-level representation is an immediate integer, physical equality is immediate equality.
- On mutable blocks at some location, or generative immutable blocks with some identity, physical equality is equality of locations or identities.
- On immutable blocks, physical-equality is under-specified, but it implies that the blocks have the same tags and their arguments are in turn physically equal.
- Two values that fall into different categories above are never physically equal.

We reflect this specification in the Zoo program logic, while keeping a reasonably high-level definition of abstract values.

Our verification work uncovered correctness bugs in existing OCaml concurrent code due to unsharing. We propose to extend the language with a [@generative] annotation for immutable constructors whose physical identity we want to reason about when implementing a fine-grained

---

[6]OCaml maintainers developing the FLAMBDA optimiser have confirmed that it may perform unsharing in certain cases.
[7]Actually, for practical reasons, we distinguish identified and unidentified generative blocks.

concurrent data structure. A rule of thumb would be to use it for all non-constant constructors involved in `Atomic`.compare_and_set operations.

6 Structural equality

Structural equality is also supported. More precisely, it is not part of the semantics of the language but implemented using low-level primitives[8]. The reason is that it is in fact difficult to specify for arbitrary values. In general, we have to compare graphs — which implies structural comparison may diverge.

Accordingly, the specification of $v_1 = v_2$ requires the (partial) ownership of a *memory footprint* corresponding to the union of the two compared graphs, giving the permission to traverse them safely. If it terminates, the comparison decides whether the two graphs are bisimilar (modulo representation conflicts, as described in Section 5). In Iris, this gives:

```
Lemma structeq_spec v1 v2 footprint,
  val_traversable footprint v1 →
  val_traversable footprint v2 →
  {{{ structeq_footprint footprint }}}
    v1 = v2
  {{{ b, RET #b;
      structeq_footprint footprint *
      ⌜(if b then val_structeq else val_structneq) footprint v1 v2⌝ }}}.
```

Obviously, this general specification is not very convenient to work with. Fortunately, for abstract values (without any mutable part), we can prove a much simpler variant saying that structural equality boils down to physical equality:

```
Lemma structeq_spec_abstract v1 v2 :
  val_abstract v1 →
  val_abstract v2 →
  {{{ True }}}
    v1 = v2
  {{{ b, RET #b; ⌜(if b then val_physeq else val_physneq) v1 v2⌝ }}}
```

This should not read as a claim that when immutable values are structurally equal then they are physically equal, but rather that they provide the exact same (structural) guarantees on those values that contain no mutable locations or generative identities.

7 Standard data structures

To save users from reinventing the wheel, we provide a library of verified standard data structures — more or less a subset of the OCaml standard library. Most of these data structures[9] are completely reimplemented in Zoo and axiom-free, including the `Array`[10] module.

Sequential data structures. We provide verified implementations of various sequential data structures: array, dynamic array (vector), list, stack, queue (bounded and unbounded), double-ended queue. We claim that the proven specifications are modular and practical. In fact, most of these data structures have already been used to verify more complex ones — we present some in Section 8 and Section 10. Especially, we developed an extensive collection of flexible specifications for the

---

[8]In OCaml, these primitives correspond to the unsafe functions `Obj`.is_int, `Obj`.tag, `Obj`.size and `Obj`.field.
[9]For practical reasons, to make them completely opaque, we chose to axiomatize a few functions from the `Domain` and `Random` modules. They could trivially be realized in Zoo.
[10]Our implementation of the `Array` module is compatible with the standard one. In particular, it uses the same low-level value representation.

$$\frac{\Psi\,1}{t.\ \text{inv}\ t\ fd\ \Psi} \qquad \frac{\begin{array}{c}\text{inv}\ t\ fd\ \Psi \\ \text{wp}\ closed\ ()\ \{\,X\ \text{false}\,\} \\ \forall\,q.\ \Psi\,q \ast\!\!-\!\! \ast \text{wp}\ open\ fd\ \{\,res.\ \Psi\,q \ast X\ \text{true}\ res\,\}\end{array}}{\begin{array}{c}\text{use}\ t\ closed\ open \\ res.\ \exists\,b.\ X\ b\ res\end{array}}$$

$$\frac{\begin{array}{c}\text{inv}\ t\ fd\ \Psi \\ \Psi\,1 \ast\!\!-\!\! \ast \exists\,chars.\ \text{unix-fd}\ fd\ 1\ chars\end{array}}{\begin{array}{c}\text{close}\ t \\ ().\ \text{closing}\ t\end{array}} \qquad \frac{\begin{array}{c}\text{inv}\ t\ fd\ \Psi \\ \text{remove}\ t\end{array}}{\begin{array}{c}\text{closing}\ t \\ o.\ \ o = \text{None} \vee o = \text{Some}\ fd \ast \Psi\,1\end{array}}$$
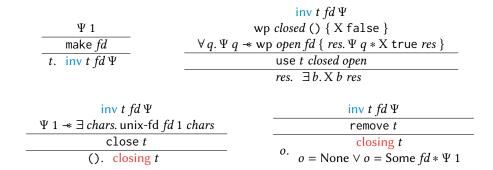
Fig. 4. Rcfd specification (excerpt)

iterators of the **Array** and **List** modules. Remarkably, our formalization of **Array** features different (fractional) predicates to express the ownership of either an entire array, a slice or even a circular slice — we used it to verify algorithms involving circular arrays, *e.g.* Chase-Lev working-stealing queue [Chase and Lev 2005].

Concurrent data structures. We provide verified implementations of various concurrent data structures: domain[11] (including domain-local storage), mutex, semaphore, condition variable, write-once variable (also known as *ivar*), atomic array. Note that there is currently no **Atomic_array** module in the OCaml standard library, but we are planning to propose it.

8 Persistent data structures

To further demonstrate the practicality of Zoo, we verified a collection of persistent data structures. This includes purely functional data structures such as persistent stack and queue, but also efficient imperative implementations of persistent array from Conchon and Filliâtre [2007], store and union-find from Allain, Clément, Moine and Scherer [2024].

Currently, verification of purely functional programs relies on the regular ZooLang translation, *i.e.* on a deeply embedded representation. However, we found this approach is cumbersome. In the future, it would be desirable to be able to verify them directly in Rocq, through a translation to Gallina. Similarly to Hacspec [Haselwarter, Hvass, Hansen, Winterhalter, Hritcu and Spitters 2024], this new translation would come with a generated proof of equivalence with the ZooLang representation.

9 Rcfd: Parallelism-safe file descriptor

As mentioned in Section 5, the **Rcfd** module from the Eio library is particularly interesting in several respects. Not only does it justify the introduction of generative constructors in OCaml, but it also demonstrates the use of Iris for expressing realistic concurrent protocols.

Specification. The **Rcfd** module provides a parallelism-safe file descriptor (FD) relying internally on reference-counting. Interestingly, it is used in Eio in two different ways, more precisely two different ownership regimes: 1) any thread can try to access or close the FD; 2) any thread can try to access the FD but only the owner thread can close it — and is responsible for closing it. To verify all uses, the specification of **Rcfd** has to support both ownership regimes. However, due to space

---

[11]Domains are the units of parallelism in OCaml 5.

constraints, we consider a simplified specification given in Figure 4. The full verified specification can be found in the mechanization. The specification features four operations[12]:

make creates a new object *t* of type **Rcfd**.t wrapping a given FD *fd*, yielding the (persistent) invariant inv *t fd* Ψ, where Ψ is an arbitrary fractional predicate. Crucially, the user must provide the full predicate Ψ 1, which is stored in the invariant. Once it is created, a wrapped FD can be accessed through the use operation and closed through the close operation.

use requires the invariant along with the weakest preconditions of the *closed* function, that is called if the FD has been flagged as closed, and *open* function, that is called if the FD is still open. To control the postconditions and the weakest preconditions, the user can choose an arbitrary predicate X parameterized by a boolean indicating whether the *closed* (false) or the *open* (true) was called. The *open* function is given a fraction of Ψ, thereby accessing the FD.

close requires the invariant and proving that the full ownership of Ψ entails the full ownership of the FD *fd*, which is necessary to call **Unix**.close. It yields closing *t*, a persistent resource witnessing that *t* has been flagged as closed. Actually, the wrapped FD is not closed immediately. It will be closed only once it is possible, meaning all ongoing calls to use owning a fraction of the FD end.

Alternatively, instead of closing the FD, remove tries to retrieve the full ownership of Ψ. To achieve it, it exploits the same mechanism as close — flagging *t* as closed as witnessed by closing *t* — but also waits until all use calls are done.

Logical state. Thomas Leonard, the author of **Rcfd**, suggested verifying it to make sure the informal concurrent protocol he described in the OCaml interface was correct. This protocol introduces a notion of monotonic logical state — modeled in Iris using a specific resource algebra [Timany and Birkedal 2021] — to describe the evolution of a FD. Originally, there were four logical states but we found that only three are necessary for the verification: **open**, **closing/users** and **closing/no-users**.

In the **open** state, the FD is available for use, meaning any thread can access it through use. Physically, this corresponds to the **Open** constructor.

When some thread flags the FD as closed through close or remove, the state transitions from **open** to **closing/users**. Crucially, there can only be one such thread. In this state, the FD is not really closed yet because of ongoing use operations. Physically, this logical transition corresponds to switching from the **Open** to the **Closing** constructor using **Atomic**.**Loc**.compare_and_set.

Once all use operations have finished, when the reference-count reaches zero, it is time to actually "close" the FD by calling the function carried by the **Closing** constructor. This has to be done only once. The "closing" thread is the one that succeeds in updating the **Closing** constructor (to a new one carrying a no-op function) using **Atomic**.**Loc**.compare_and_set. At this point, the state transitions from **closing/users** to **closing/no-users** and the wrapper no longer owns the FD.

Generative contructors. As explained in Section 5, the **Open** constructor has to be generative to prevent *unsharing*. In fact, the **Closing** constructor also has to be generative to prevent *sharing*, otherwise two calls could have a shared value of next and believe they both won the second update.

10 Saturn: A library of standard lock-free data structures

We verified a collection of standard lock-free data structures from the Saturn, Eio and Picos [Karvonen 2025b] libraries. It includes stacks, queues (list-based, array-based and stack-based) and bags. These data structures are meant to be used as is or adapted to fit specific needs. To cover a wide range of use cases, we provide specialized variants: bounded or unbounded, single-producer (SP) or multi-producer (MP), single-consumer (SC) or multi-consumer (MC).

---

[12]We omitted two non-essential operations: is_open and peek.

Due to space constraints, we focus on the most important algorithms and refrain from showing the corresponding (non-trivial) IRIS invariants, which are mechanized in ROCQ.

10.1 Stacks

We verified three variants of the Treiber stack [Treiber 1986]: 1) unbounded MPMC (the standard one), 2) bounded MPMC, 3) closable unbounded MPMC. This last variant features a closing mechanism: at some point, some thread can decide to close the stack, retrieving the current content and preventing others from operating on it. For example, we used it to represent a set of vertex successors in the context of a concurrent graph implementation (not presented in this paper).

As explained in Section 5, the three verified stacks use generative constructors to prevent sharing. One may ask whether it would be easier to use a mutable version of lists instead. From the programmer's perspective, this is unsatisfactory because 1) the compiler will typically emit warnings complaining that the mutability is not exploited and 2) it does not really reflect the intent, *i.e.* we want precise guarantees for physical equality, not modify the list. From the verification perspective, this is also unsatisfactory because the mutable representation is more complex to write and reason about: pointers and points-to assertions versus pure ROCQ list.

Although verified stacks may seem like a not-so-new contribution, it is, as far as we know, the first verification of realistic OCAML implementations. For comparison, the exemplary concurrent stacks verified in IRIS [Iris development team 2025] all suffer from the same flaw: they need to introduce indirections (pointers) to be able to use the compare-and-set primitive.

10.2 List-based queues

We verified four variants of the Michael-Scott queue [Michael and Scott 1996]: unbounded MPMC (the standard one), bounded MPMC, unbounded MPSC and unbounded SPMC. The SPMC queue is used to implement one of the bags — a relaxed queue guaranteeing only per-producer ordering.

In the IRIS literature, Vindum and Birkedal [2021] established contextual refinement of the Michael-Scott queue while Mulder and Krebbers [2023] proved logical atomicity. However, we had to redesign and extend the invariant for several reasons.

Efficient implementation. The Michael-Scott essentially consists of a singly linked list of nodes that only grows over time. The previously verified implementations, implemented in HEAPLANG, use a double indirection to represent the list [Vindum and Birkedal 2021, Figure 2]. Similarly to the Treiber stack, this is made so as to be able to use the compare-and-set primitive of HEAPLANG. Vindum and Birkedal [2021] write:

> A node is a pointer to either none or some of a pair of a value and a pointer to
> the next node. The pointer serves to make nodes comparable by pointer equality
> such that pointers to nodes can be changed with CAS.

In OCaml, this would correspond to introducing extra atomic references (`Atomic.t`) between the nodes. Using atomic record fields, we can represent the list more efficiently, without the extra indirection. However, there is one subtlety: in this new representation, we need to clear the outdated nodes so that their value is no longer reachable and can be garbage-collected, in other words to prevent a memory leak. This subtlety is not discussed in the original implementation [Michael and Scott 1996] designed for non-garbage-collected languages, but it is folklore; it is implemented in Saturn (saturn#64) and we verified that it preserves correctness.

To deal with this representation in separation logic, we introduce the notion of *explicit chain* that allows decoupling the chain structure formed by the nodes and the content of the nodes. Concretely, the assertion xchain dq ℓs dst represents a chain linking locations ℓs and ending at value dst; dq is a discardable fraction [Vindum and Birkedal 2021] that controls the ownership of the chain. This

notion is very flexible as it is independent of the rest of the structure. As a matter of fact, we used it and its generalization to doubly linked list more broadly, to verify other algorithms. All the variants of Michael-Scott we verified rely on it. In particular, it was quite straightforward to extend the invariant of the bounded queue, where nodes carry more (mutable and immutable) information.

External linearization point. Our work also revealed another interesting aspect that is not addressed in the literature, as far as we know. None of the previously verified implementations deal with the `is_empty` operation, that consists in reading the sentinel node and checking whether it has a successor. It it has no successor, it is necessarily the last node of the chain, hence the queue is empty. If it does have a successor, `is_empty` returns `false`, meaning we must have observed a non-empty queue. However, this last part is more tricky than it may seem. Indeed, it may happen that 1) we read the sentinel while the queue is empty, 2) other operations fill and empty again the queue so that the sentinel is outdated, 3) we read the successor of the former sentinel while the queue is still empty. The crucial point here is that `is_empty` is linearized when the first push operation filled the queue. In other words, the linearization point of `is_empty` is triggered by another operation; this is called an *external linearization point*. To handle this in the proof, we introduced a mechanism in the invariant to transfer the IRIS resource materializing the linearization point[13] from `is_empty` to push and vice versa.

10.3   Stack-based queues

A standard way to implement a sequential queue is to use two stacks: producers push onto the *back stack* while consumers pop from the *front stack*, stealing and reversing the back stack when needed. Based on this simple idea, Vesa Karvonen developed a new lock-free concurrent queue. We verified the MPMC variant used in `Picos` and the closable MPSC variant used in `Eio`.

As in the sequential implementation, the two stacks are mainly immutable. Both stacks are updated using compare-and-set so we use generative constructors to reason about physical equality.

Similarly again, producers and consumers work concurrently on separate stacks, limiting interference. The key difference compared to the sequential version is that the algorithm has to deal with the concurrent back stack reversal in a lock-free manner. Essentially, the concurrent protocol — and therefore the IRIS invariant — includes a *destabilization* phase during which a new back stack pointing to the former one awaits to be *stabilized*, which happens when the reversed former back stack becomes the new front stack. In practice, the synchronization is a bit tricky and relies on the indices of the elements.

11   Memory safety

Concurrency creates tensions between performance and memory-safety. The OCaml maintainers intend to maintain memory-safety for all OCaml programs, including racy concurrent programs. They have tried to imprint this focus on memory-safety to library authors as well.

Performance-sensitive OCaml libraries are sometimes written using unsafe primitives, that may break memory-safety if used incorrectly. It is the responsibility of code authors to ensure that the preconditions of those unsafe primitives are satisfied to ensure safety. Unfortunately, the addition to parallel code execution in OCaml 5 broke the safety of some existing code: it adds more possible interleaving and may invalidate safety reasoning.

---

[13]This resource is known as an *atomic update*. Mulder and Krebbers [2023] provide a good description.

11.1 Dynarray

In January 2023, Gabriel Scherer proposed the addition of **Dynarray**, a module of (sequential) resizable arrays, to the OCaml standard libary.[14]

A concurrent safety problem. Resizable arrays are implemented as a record with two mutable fields, a size field that stores the current length of the resizable array, and a data field storing the "backing array" a (non-resizable) array of size at least size elements. When a user adds a new element to the end of a resizable array, it typically suffices to write the new element at index size in the backing array, and then increment the size field. But when the size field reaches the actual length of the backing array (which we call the "capacity" of the resizable array), we first need to allocate a new, larger backing array, to copy the values from the old to the new backing array, and to overwrite the data field with the new backing array.

This operation of adding a new element is performance-critical for resizable arrays. After the size check and potential resizing, sequentially we know that the backing array has enough space, and we can write the new element using an unsafe_set primitive that avoids a redundant bound-check on array access. This optimization can provide speedups of up to 20% in certain scenarios.

```
let add_last a x = (* sequential version *)
  let size = a.size in
  if Array.length a.data = size then ensure_capacity a (size + 1);
  a.size <- size + 1;
  Array.unsafe_set a.data size x
```

Unfortunately, the safety reasoning becomes incorrect in presence of parallelism: another thread could mutate the backing array after the size check and before the unsafe write, for example **Dynarray**.reset which sets the size to 0 and replaces the backing array by an empty array. Writing to a (non-resizable) array outside its bounds without bound checks breaks memory safety, so this implementation was memory-safe under OCaml 4 but it becomes memory-unsafe under OCaml 5. On the other hand, **Dynarray** is explicitly documented as a sequential data structure, so it is the user responsibility to respect this precondition by using appropriate synchronization (for example a mutex) to prevent data races. Some performance-obsessed users would argue that if *other* users fail in their responsibility of ensuring sequential access, then they do not deserve memory-safety. The OCaml maintainers and standard library authors consider on the other hand that memory-safety should be preserved even in this case: it can only be lifted for operations that are explicitly marked as unsafe, and hopefully have simple, checkable preconditions. So they decided to change the optimization of **Dynarray** to guarantee memory-safety even in presence of racy concurrent usage.

The proposed implementation reads the data field to get the backing array b and performs the resizing check. If no resizing is necessary, then an unsafe write is performed as before (on the backing array b, without re-reading the data field). In the infrequent case where resizing is necessary, the operation retries afterwards. (A bound-checking write would suffice for safety.)

```
let rec add_last a x = (* memory-safe under concurrency *)
  let {size; data} = a in
  if Array.length data >= size
  then ( ensure_capacity a (size + 1); add_last a x )
  else ( a.size <- size + 1; Array.unsafe_set arr size x )
```

A concurrent verification problem. We verified (a representative fragment of) the proposed **Dynarray** implementation in ZooLang, proving that it is functionally correct under sequential

---

[14]https://github.com/ocaml/ocaml/pull/11882

usage, and that it does preserve memory-safety even under concurrent usage. The verified fragment is now part of the Zoo standard library, and can be used in further verification projects.

Informally, the verification relies on two different invariants:

- Functional correctness relies on a *strong* invariant, which may not be preserved under concurrent usage; for example that the content of the size field is smaller than the size of the backing array.
- Memory-safety relies on a *weak* invariant that does not suffice to prove correctness, but is preserved under concurrent usage. For this **Dynarray** implementation, the weak invariant is that the backing array remains well-typed and that the content of the size field is non-negative.

To informally check the two desired properties, it suffices to check that the strong invariant is indeed a sequential invariant, and that implies functional correctness; and separately that the weak invariant is preserved by all operations on the data structure.

In our mechanized verification, the strong invariant becomes a *model* of the data structure, and we verify separation-logic triples where each operation can assume unique ownership of the model in input, and has to return ownership of a valid model in output, as is standard. For example, the "functional correctness" lemma for the add_last function is the following:

```
Lemma dynarray_2_push_spec t vs v :
  {{{ dynarray_2_model t vs }}}
    dynarray_2_push t v
  {{{ RET (); dynarray_2_model t (vs ++ [v]) }}}.
```

This lemma establishes correctness in a purely-sequential usage, but also under any concurrent usage where correct synchronization is used to guarantee unique ownership of the model.

The formal counterpart of our weak invariant is a *semantic type*, following the approach of the RustBelt project [Jung, Jourdan, Krebbers and Dreyer 2018], which also deals with unsafe fragments within a language intended to be safe. Note that the RustBelt semantic types are derived from Rust types, where mutating operations get a mutable borrow and thus unique ownership (for a time). In contrast, our semantic types carry no ownership of the values they manipulate, so they are stored entirely as invariants, and all interactions with the structure must be shown to preserve this invariant atomically. The invariant must be robust against any interference coming from any other function called in parallel on the same structure.

The definition of the semantic type for dynarrays, and the statement of memory safety (or in fact semantic typing) for add_last look as follows:

```
Definition itype_dynarray_2 ty t : iProp Σ :=
  ∃ l,
  ⌜t = #l⌝ *
  inv nroot (
    ∃ (sz : nat) cap data,
    l.[size] ↦ #sz *
    l.[data] ↦ data * itype_array ty cap data
  ).

Lemma dynarray_2_push_type ty t v :
  {{{ itype_dynarray_2 ty t * ty v }}}
    dynarray_2_push ty t v
  {{{ RET (); True }}}.
```

The predicate ty is the semantic type for the elements of the array — it must hold for each value of the backing array and for the new value v.

11.2 Saturn: Single-producer or Single-consumer queues

Similar problems of memory-safety occur in the concurrent data structures of the Saturn library. For example, we explained (Section 10.2) that the Saturn implementation of the Michael-Scott MPMC queue is careful to "erase" values stored in the queue to avoid memory leaks. This erasure is performed by writing a non-type-safe dummy value, **Obj**.magic (). We define a semantic type for the MPMC queue which is essentially its concurrent invariant, and prove that this unsafe idiom does not endanger memory-safety.

On the other hand, the efficient implementations of single-consumer or single-producer queues in Saturn do *not* respect the general OCaml recommendation of favoring safety over performance: if a library user uses a single-consumer structure from two consumers in racy ways, they lose memory safety. Our correctness results imply memory-safety when the single-consumer or single-producer protocol is respected; when the caller does own its end of the structure uniquely, as our precondition requires. But we cannot prove memory-safety with only persistent preconditions, it does not hold.

The Saturn authors also provide "safe" variants of their data structures, which are slightly slower but memory-safe even for unintended concurrent usage. This typically adds an indirection, such as using the type 'a option instead of 'a, which provides a type-safe **None** value for dummies.

12 Related work

We have amply covered related work on verification of concurrent fine-grained data structures in the main body of the paper. We focus here on the related work on program verification and the question of physical equality.

12.1 Non-automated verification

The verified program is translated, manually or in an automated way, into a representation living inside a proof assistant. The user has to write specifications and prove them.

Translating into the native language of the proof assistant, such as Gallina for Rocq, is challenging as it is hard to faithfully preserve the semantics of the source language, which typically has non-terminating functions for example. Monadic translations should support it, but faithfully encoding all impure behaviors is challenging, and tools typically provide a best-effort translation [Claret 2025; Spector-Zabusky, Breitner, Rizkallah and Weirich 2018] that is only approximately sound.

The representation may be embedded, meaning the semantics of the language is formalized in the proof assistant. This is the path taken by some recent works (for example Gondelman, Hinrichsen, Pereira, Timany and Birkedal [2023]) harnessing the power of separation logic. In particular, CFML [Charguéraud 2023] and Osiris [Daby-Seesaram, Madiot, Pottier, Seassau and Yoon 2024] target OCaml. However, CFML does not support concurrency and is not based on Iris. Osiris, still under development, is based on Iris but does not support concurrency.

Prior to our work, HeapLang was thus the most appropriate tool to verify concurrent OCaml programs. We discussed limitations of HeapLang in the introduction, and ZooLang is our proposal to improve on this. Conversely, one notable limitation of ZooLang today is its lack of support for OCaml's relaxed memory model.

12.2 Semi-automated verification

In semi-automated verification approaches, the verified program is annotated by the user to guide the verification tool: preconditions, postconditions, invariants, *etc*. Given this input, the

verification tool generates proof obligations that are mostly automatically discharged. One may further distinguish two types of semi-automated systems: *foundational* and *non-foundational*.

In *non-foundational* automated verification, the tool and external solvers it may rely on are part of the trusted computing base. It is the most common approach and has been widely applied in the literature [Astrauskas, Bilý, Fiala, Grannan, Matheja, Müller, Poli and Summers 2022; Denis, Jourdan and Marché 2022; Filliâtre and Paskevich 2013; Jacobs, Smans, Philippaerts, Vogels, Penninckx and Piessens 2011; Lattuada, Hance, Cho, Brun, Subasinghe, Zhou, Howell, Parno and Hawblitzel 2023; Müller, Schwerhoff and Summers 2017; Pulte, Makwana, Sewell, Memarian, Sewell and Krishnaswami 2023; Swamy, Chen, Fournet, Strub, Bhargavan and Yang 2013], including to OCaml by Cameleer [Pereira and Ravara 2021], which uses the Gospel specification language [Charguéraud, Filliâtre, Lourenço and Pereira 2019] and Why3 [Filliâtre and Paskevich 2013].

In *foundational* automated verification, proofs are checked by a proof assistant so the automation does not have to be trusted. To our knowledge, it has been applied to C [Sammler, Lepigre, Krebbers, Memarian, Dreyer and Garg 2021] and Rust [Gäher, Sammler, Jung, Krebbers and Dreyer 2024].

Zoo is mostly non-automated — except for our use Diaframe for local automation of separation logic reasoning. We would be interested in moving towards more automation in the future.

12.3 Physical equality

There is some literature in proof-assistant research on reflecting physical equality from the implementation language into the proof assistant, for optimization purposes: for example, exposing OCaml's physical equality as a predicate in Rocq lets us implement some memoization and sharing techniques in Rocq libraries. However, axiomatizing physical equality in the proof assistant is difficult, and can result in inconsistencies.

The earlier discussions of this question that we know come from Jourdan's thesis [Jourdan 2016] (chapter 9), also presented more succintly in [Braibant, Jourdan and Monniaux 2014]. This work introduces the Jourdan condition, that physical equality implies equality of values. [Boulmé 2021] extends the treatment of physical equality in Rocq, integrating it in an "extraction monad" to control it more safely. There is also a discussion of similar optimizations in Lean in [Selsam, Hudon and de Moura 2020].

The correctness of the axiomatization of physical equality depends on the type of the values being compared: axiomatizations are typically polymorphic on any type A, but their correctness depends on the specific A being considered. For example, it is easy to correctly characterize physical equality on natural numbers, and other non-dependent types arising in Rocq verification projects. One difficulty in HeapLang and ZooLang is that they are untyped languages, their representation of `0` and `false` has the same type. But our remark that structural equality (in OCaml) does not necessarily coincide with definitional equality (in Rocq) also applies to other Rocq types: our examples with an existential **Any** constructor (see Section 5) can be reproduced with Σ-types.

13 Conclusion and future work

We presented Zoo, a framework for verification of concurrent OCaml 5 programs. While it is not yet available on opam, it can be installed and used in other Rocq projects. Zoo has been used to verify sequential imperative algorithms and a significant library of lock-free data structures. Its main weakness so far is its memory model, which is sequentially consistent as opposed to the relaxed OCaml 5 memory model. In the future we also plan to add exceptions and algebraic effects.

Another interesting direction would be to combine Zoo with semi-automated techniques. Similarly to Why3, the simple parts of the verification effort would be done in a semi-automated way, while the most difficult parts would be conducted in Rocq.

References

Clément Allain, Basile Clément, Alexandre Moine, and Gabriel Scherer. 2024. Snapshottable Stores. *Proc. ACM Program. Lang.* 8, ICFP (2024), 338–369. doi:10.1145/3674637

Vytautas Astrauskas, Aurel Bilý, Jonás Fiala, Zachary Grannan, Christoph Matheja, Peter Müller, Federico Poli, and Alexander J. Summers. 2022. The Prusti Project: Formal Verification for Rust. In *NASA Formal Methods - 14th International Symposium, NFM 2022, Pasadena, CA, USA, May 24-27, 2022, Proceedings (Lecture Notes in Computer Science, Vol. 13260)*, Jyotirmoy V. Deshmukh, Klaus Havelund, and Ivan Perez (Eds.). Springer, 88–108. doi:10.1007/978-3-031-06773-0_5

Lars Birkedal, Thomas Dinsdale-Young, Armaël Guéneau, Guilhem Jaber, Kasper Svendsen, and Nikos Tzevelekos. 2021. Theorems for free from separation logic specifications. *Proc. ACM Program. Lang.* 5, ICFP (2021), 1–29. doi:10.1145/3473586

Sylvain Boulmé. 2021. *Formally Verified Defensive Programming (efficient Coq-verified computations from untrusted ML oracles)*. Accreditation to supervise research. Université Grenoble-Alpes. https://hal.science/tel-03356701 See also http://www-verimag.imag.fr/ boulme/hdr.html.

Thomas Braibant, Jacques-Henri Jourdan, and David Monniaux. 2014. Implementing and Reasoning About Hash-consed Data Structures in Coq. *J. Autom. Reason.* 53, 3 (2014), 271–304. doi:10.1007/S10817-014-9306-0

Tej Chajed, Joseph Tassarotti, M. Frans Kaashoek, and Nickolai Zeldovich. 2019. Verifying concurrent, crash-safe systems with Perennial. In *Proceedings of the 27th ACM Symposium on Operating Systems Principles, SOSP 2019, Huntsville, ON, Canada, October 27-30, 2019*, Tim Brecht and Carey Williamson (Eds.). ACM, 243–258. doi:10.1145/3341301.3359632

Yun-Sheng Chang, Ralf Jung, Upamanyu Sharma, Joseph Tassarotti, M. Frans Kaashoek, and Nickolai Zeldovich. 2023. Verifying vMVCC, a high-performance transaction library using multi-version concurrency control. In *17th USENIX Symposium on Operating Systems Design and Implementation, OSDI 2023, Boston, MA, USA, July 10-12, 2023*, Roxana Geambasu and Ed Nightingale (Eds.). USENIX Association, 871–886. https://www.usenix.org/conference/osdi23/presentation/chang

Arthur Charguéraud. 2023. *Habilitation thesis: A Modern Eye on Separation Logic for Sequential Programs. (Un nouveau regard sur la Logique de Séparation pour les programmes séquentiels)*. Université de Strasbourg. https://tel.archives-ouvertes.fr/tel-04076725

Arthur Charguéraud, Jean-Christophe Filliâtre, Cláudio Lourenço, and Mário Pereira. 2019. GOSPEL - Providing OCaml with a Formal Specification Language. In *Formal Methods - The Next 30 Years - Third World Congress, FM 2019, Porto, Portugal, October 7-11, 2019, Proceedings (Lecture Notes in Computer Science, Vol. 11800)*, Maurice H. ter Beek, Annabelle McIver, and José N. Oliveira (Eds.). Springer, 484–501. doi:10.1007/978-3-030-30942-8_29

David Chase and Yossi Lev. 2005. Dynamic circular work-stealing deque. In *SPAA 2005: Proceedings of the 17th Annual ACM Symposium on Parallelism in Algorithms and Architectures, July 18-20, 2005, Las Vegas, Nevada, USA*, Phillip B. Gibbons and Paul G. Spirakis (Eds.). ACM, 21–28. doi:10.1145/1073970.1073974

Guillaume Claret. 2025. coq-of-ocaml. https://github.com/formal-land/coq-of-ocaml

Sylvain Conchon and Jean-Christophe Filliâtre. 2007. A persistent union-find data structure. In *Proceedings of the ACM Workshop on ML, 2007, Freiburg, Germany, October 5, 2007*, Claudio V. Russo and Derek Dreyer (Eds.). ACM, 37–46. doi:10.1145/1292535.1292541

Pedro da Rocha Pinto, Thomas Dinsdale-Young, and Philippa Gardner. 2014. TaDA: A Logic for Time and Data Abstraction. In *ECOOP 2014 - Object-Oriented Programming - 28th European Conference, Uppsala, Sweden, July 28 - August 1, 2014. Proceedings (Lecture Notes in Computer Science, Vol. 8586)*, Richard E. Jones (Ed.). Springer, 207–231. doi:10.1007/978-3-662-44202-9_9

Arnaud Daby-Seesaram, Jean-Marie Madiot, François Pottier, Remy Seassau, and Irene Yoon. 2024. Osiris. https://gitlab.inria.fr/fpottier/osiris

Xavier Denis, Jacques-Henri Jourdan, and Claude Marché. 2022. Creusot: A Foundry for the Deductive Verification of Rust Programs. In *Formal Methods and Software Engineering - 23rd International Conference on Formal Engineering Methods, ICFEM 2022, Madrid, Spain, October 24-27, 2022, Proceedings (Lecture Notes in Computer Science, Vol. 13478)*, Adrián Riesco and Min Zhang (Eds.). Springer, 90–105. doi:10.1007/978-3-031-17244-1_6

Stephen Dolan, KC Sivaramakrishnan, and Anil Madhavapeddy. 2018. Bounding data races in space and time. In *PLDI*.

Jean-Christophe Filliâtre and Andrei Paskevich. 2013. Why3 - Where Programs Meet Provers. In *Programming Languages and Systems - 22nd European Symposium on Programming, ESOP 2013, Held as Part of the European Joint Conferences on Theory and Practice of Software, ETAPS 2013, Rome, Italy, March 16-24, 2013. Proceedings (Lecture Notes in Computer Science, Vol. 7792)*, Matthias Felleisen and Philippa Gardner (Eds.). Springer, 125–128. doi:10.1007/978-3-642-37036-6_8

Lennard Gäher, Michael Sammler, Ralf Jung, Robbert Krebbers, and Derek Dreyer. 2024. RefinedRust: A Type System for High-Assurance Verification of Rust Programs. *Proc. ACM Program. Lang.* 8, PLDI (2024), 1115–1139. doi:10.1145/3656422

Léon Gondelman, Jonas Kastberg Hinrichsen, Mário Pereira, Amin Timany, and Lars Birkedal. 2023. Verifying Reliable Network Components in a Distributed Separation Logic with Dependent Separation Protocols. *Proc. ACM Program. Lang.* 7, ICFP (2023), 847–877. doi:10.1145/3607859

Philipp G. Haselwarter, Benjamin Salling Hvass, Lasse Letager Hansen, Théo Winterhalter, Catalin Hritcu, and Bas Spitters. 2024. The Last Yard: Foundational End-to-End Verification of High-Speed Cryptography. In *Proceedings of the 13th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2024, London, UK, January 15-16, 2024*, Amin Timany, Dmitriy Traytel, Brigitte Pientka, and Sandrine Blazy (Eds.). ACM, 30–44. doi:10.1145/3636501.3636961

Maurice Herlihy and Jeannette M. Wing. 1990. Linearizability: A Correctness Condition for Concurrent Objects. *ACM Trans. Program. Lang. Syst.* 12, 3 (1990), 463–492. doi:10.1145/78969.78972

Iris development team. 2025. Iris examples. https://gitlab.mpi-sws.org/iris/examples/

Bart Jacobs, Jan Smans, Pieter Philippaerts, Frédéric Vogels, Willem Penninckx, and Frank Piessens. 2011. VeriFast: A Powerful, Sound, Predictable, Fast Verifier for C and Java. In *NASA Formal Methods - Third International Symposium, NFM 2011, Pasadena, CA, USA, April 18-20, 2011. Proceedings (Lecture Notes in Computer Science, Vol. 6617)*, Mihaela Gheorghiu Bobaru, Klaus Havelund, Gerard J. Holzmann, and Rajeev Joshi (Eds.). Springer, 41–55. doi:10.1007/978-3-642-20398-5_4

Jacques-Henri Jourdan. 2016. *Verasco: a Formally Verified C Static Analyzer. (Verasco: un analyseur statique pour C formellement vérifié).* Ph. D. Dissertation. Paris Diderot University, France. https://tel.archives-ouvertes.fr/tel-01327023

Ralf Jung, Jacques-Henri Jourdan, Robbert Krebbers, and Derek Dreyer. 2018. RustBelt: securing the foundations of the rust programming language. *PACMPL* 2, POPL (2018). doi:10.1145/3158154

Ralf Jung, Robbert Krebbers, Jacques-Henri Jourdan, Ales Bizjak, Lars Birkedal, and Derek Dreyer. 2018. Iris from the ground up: A modular foundation for higher-order concurrent separation logic. *J. Funct. Program.* 28 (2018), e20. doi:10.1017/S0956796818000151

Ralf Jung, Rodolphe Lepigre, Gaurav Parthasarathy, Marianna Rapoport, Amin Timany, Derek Dreyer, and Bart Jacobs. 2020. The future is ours: prophecy variables in separation logic. *Proc. ACM Program. Lang.* 4, POPL (2020), 45:1–45:32. doi:10.1145/3371113

Vesa Karvonen. 2025a. Kcas. https://github.com/ocaml-multicore/kcas

Vesa Karvonen. 2025b. Picos. https://github.com/ocaml-multicore/picos

Vesa Karvonen and Carine Morel. 2025. Saturn. https://github.com/ocaml-multicore/saturn

Robbert Krebbers, Jacques-Henri Jourdan, Ralf Jung, Joseph Tassarotti, Jan-Oliver Kaiser, Amin Timany, Arthur Charguéraud, and Derek Dreyer. 2018. MoSeL: a general, extensible modal framework for interactive proofs in separation logic. *Proc. ACM Program. Lang.* 2, ICFP (2018), 77:1–77:30. doi:10.1145/3236772

Andrea Lattuada, Travis Hance, Chanhee Cho, Matthias Brun, Isitha Subasinghe, Yi Zhou, Jon Howell, Bryan Parno, and Chris Hawblitzel. 2023. Verus: Verifying Rust Programs using Linear Ghost Types. *Proc. ACM Program. Lang.* 7, OOPSLA1 (2023), 286–315. doi:10.1145/3586037

Anil Madhavapeddy and Thomas Leonard. 2025. Eio. https://github.com/ocaml-multicore/eio

Glen Mével, Jacques-Henri Jourdan, and François Pottier. 2020. Cosmo: a concurrent separation logic for multicore OCaml. *Proc. ACM Program. Lang.* 4, ICFP (2020), 96:1–96:29. doi:10.1145/3408978

Maged M. Michael and Michael L. Scott. 1996. Simple, Fast, and Practical Non-Blocking and Blocking Concurrent Queue Algorithms. In *Proceedings of the Fifteenth Annual ACM Symposium on Principles of Distributed Computing, Philadelphia, Pennsylvania, USA, May 23-26, 1996*, James E. Burns and Yoram Moses (Eds.). ACM, 267–275. doi:10.1145/248052.248106

Ike Mulder and Robbert Krebbers. 2023. Proof Automation for Linearizability in Separation Logic. *Proc. ACM Program. Lang.* 7, OOPSLA1 (2023), 462–491. doi:10.1145/3586043

Ike Mulder, Robbert Krebbers, and Herman Geuvers. 2022. Diaframe: automated verification of fine-grained concurrent programs in Iris. In *PLDI '22: 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation, San Diego, CA, USA, June 13 - 17, 2022*, Ranjit Jhala and Isil Dillig (Eds.). ACM, 809–824. doi:10.1145/3519939.3523432

Peter Müller, Malte Schwerhoff, and Alexander J. Summers. 2017. Viper: A Verification Infrastructure for Permission-Based Reasoning. In *Dependable Software Systems Engineering*, Alexander Pretschner, Doron Peled, and Thomas Hutzelmann (Eds.). NATO Science for Peace and Security Series - D: Information and Communication Security, Vol. 50. IOS Press, 104–125. doi:10.3233/978-1-61499-810-5-104

Mário Pereira and António Ravara. 2021. Cameleer: A Deductive Verification Tool for OCaml. In *Computer Aided Verification - 33rd International Conference, CAV 2021, Virtual Event, July 20-23, 2021, Proceedings, Part II (Lecture Notes in Computer Science, Vol. 12760)*, Alexandra Silva and K. Rustan M. Leino (Eds.). Springer, 677–689. doi:10.1007/978-3-030-81688-9_31

Christopher Pulte, Dhruv C. Makwana, Thomas Sewell, Kayvan Memarian, Peter Sewell, and Neel Krishnaswami. 2023. CN: Verifying Systems C Code with Separation-Logic Refinement Types. *Proc. ACM Program. Lang.* 7, POPL (2023), 1–32. doi:10.1145/3571194

Michael Sammler, Rodolphe Lepigre, Robbert Krebbers, Kayvan Memarian, Derek Dreyer, and Deepak Garg. 2021. RefinedC: automating the foundational verification of C code with refined ownership types. In *PLDI '21: 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation, Virtual Event, Canada, June 20-25, 2021*, Stephen N. Freund and Eran Yahav (Eds.). ACM, 158–174. doi:10.1145/3453483.3454036

Daniel Selsam, Simon Hudon, and Leonardo de Moura. 2020. Sealing pointer-based optimizations behind pure functions. *Proc. ACM Program. Lang.* 4, ICFP, Article 115 (Aug. 2020), 20 pages. doi:10.1145/3408997

KC Sivaramakrishnan, Stephen Dolan, Leo White, Sadiq Jaffer, Tom Kelly, Anmol Sahoo, Sudha Parimala, Atul Dhiman, and Anil Madhavapeddy. 2020. Retrofitting parallelism onto OCaml. *Proc. ACM Program. Lang.* 4, ICFP, Article 113 (Aug. 2020), 30 pages. doi:10.1145/3408995

Antal Spector-Zabusky, Joachim Breitner, Christine Rizkallah, and Stephanie Weirich. 2018. Total Haskell is reasonable Coq. In *Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los Angeles, CA, USA, January 8-9, 2018*, June Andronick and Amy P. Felty (Eds.). ACM, 14–27. doi:10.1145/3167092

Nikhil Swamy, Juan Chen, Cédric Fournet, Pierre-Yves Strub, Karthikeyan Bhargavan, and Jean Yang. 2013. Secure distributed programming with value-dependent types. *J. Funct. Program.* 23, 4 (2013), 402–451. doi:10.1017/S0956796813000142

Amin Timany and Lars Birkedal. 2021. Reasoning about monotonicity in separation logic. In *CPP '21: 10th ACM SIGPLAN International Conference on Certified Programs and Proofs, Virtual Event, Denmark, January 17-19, 2021*, Catalin Hritcu and Andrei Popescu (Eds.). ACM, 91–104. doi:10.1145/3437992.3439931

R. K. Treiber. 1986. *Systems Programming: Coping with Parallelism.* International Business Machines Incorporated, Thomas J. Watson Research Center. https://books.google.fr/books?id=YQg3HAAACAAJ

Simon Friis Vindum and Lars Birkedal. 2021. Contextual refinement of the Michael-Scott queue (proof pearl). In *CPP '21: 10th ACM SIGPLAN International Conference on Certified Programs and Proofs, Virtual Event, Denmark, January 17-19, 2021*, Catalin Hritcu and Andrei Popescu (Eds.). ACM, 76–90. doi:10.1145/3437992.3439930

Simon Friis Vindum, Dan Frumin, and Lars Birkedal. 2022. Mechanized verification of a fine-grained concurrent queue from meta's folly library. In *CPP '22: 11th ACM SIGPLAN International Conference on Certified Programs and Proofs, Philadelphia, PA, USA, January 17 - 18, 2022*, Andrei Popescu and Steve Zdancewic (Eds.). ACM, 100–115. doi:10.1145/3497775.3503689